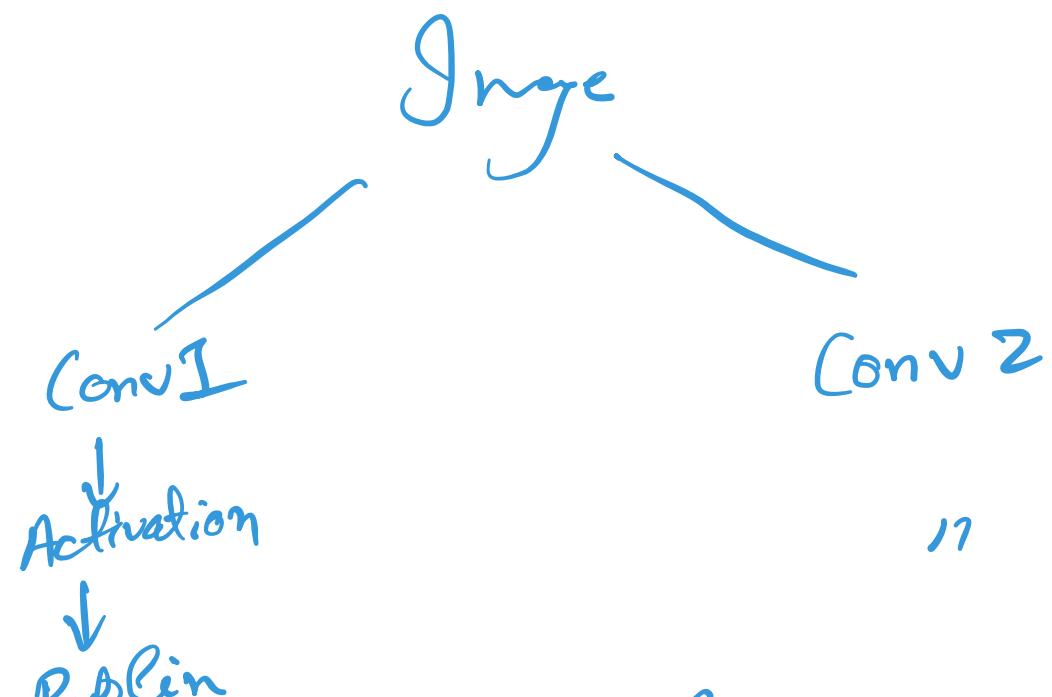
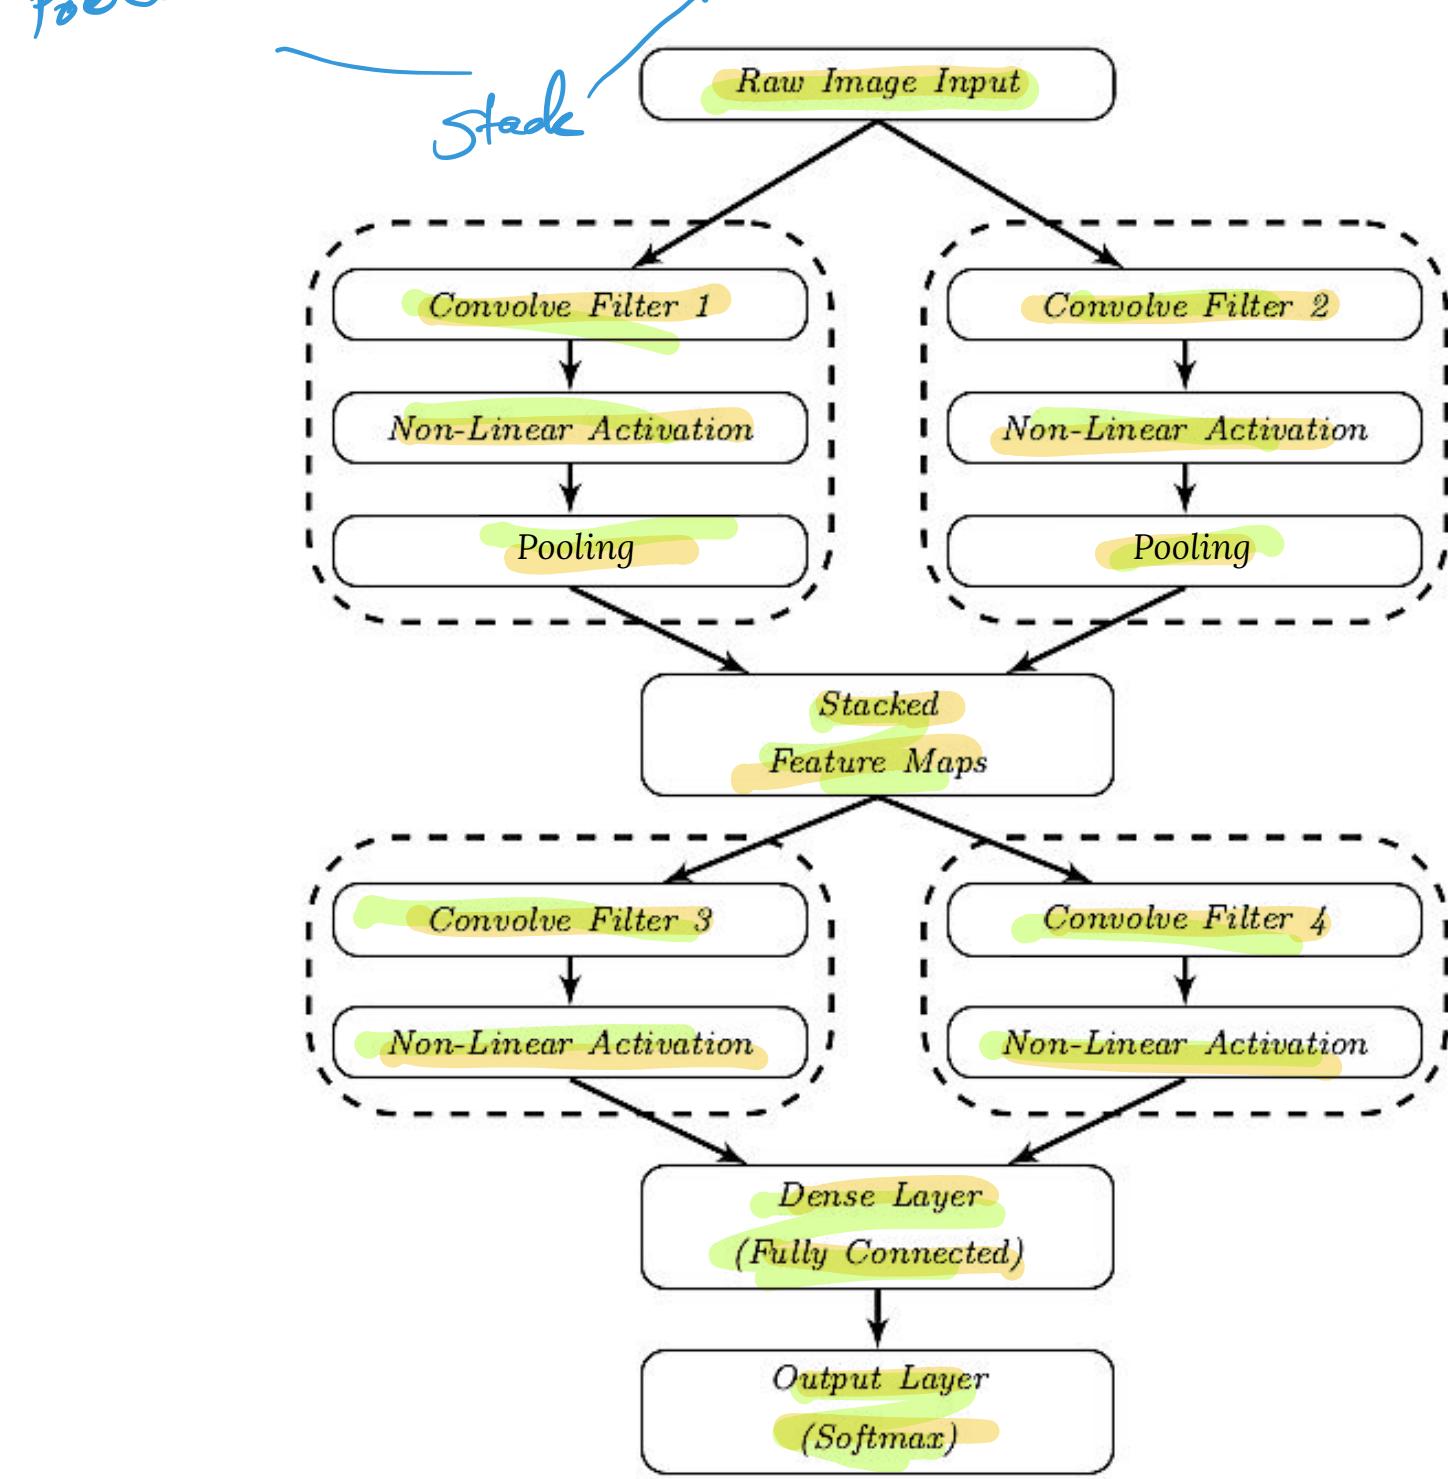


# Deep learning for Computer Vision



# Convolution Neural Networks



Schematic of typical sequences of layers in CNNs

# Special Convolution (Depth-wise separable convolutions)

Ek single channel layer ka  
Conv. layer fir  $1 \times 1$  se overall  
filters bache hain i.e Depthwise

# Separable convolutions.

## Depth-wise separable convolutions

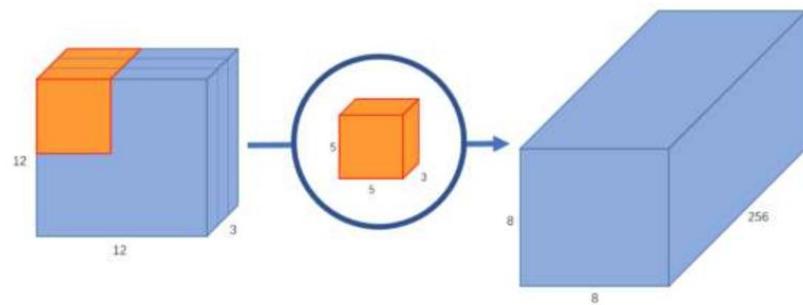
```
from keras.layers import Input, DepthwiseConv2D, Conv2D, BatchNormalization, Activation

# Input Layer for an image of size 12x12 with 3 channels
inputs = Input(shape=(12, 12, 3))

# Depthwise convolutional layer
depthwise_conv = DepthwiseConv2D(kernel_size=(5, 5), padding='valid')(inputs)
depthwise_conv = Activation('relu')(depthwise_conv)

# Pointwise convolutional layer
pointwise_conv = Conv2D(filters=256, kernel_size=(1, 1))(depthwise_conv)
pointwise_conv = Activation('relu')(pointwise_conv)
```

But why? → reductions in  
params as well as multiplying



- Each filter has a kernel size of  $5 \times 5 \times 3$  (there are 3 input channels). Plus one bias term for each filter.
- Total number of parameters in the Convolutional layer is  $256 \times (5 \times 5 \times 3 + 1) = 19456$ .

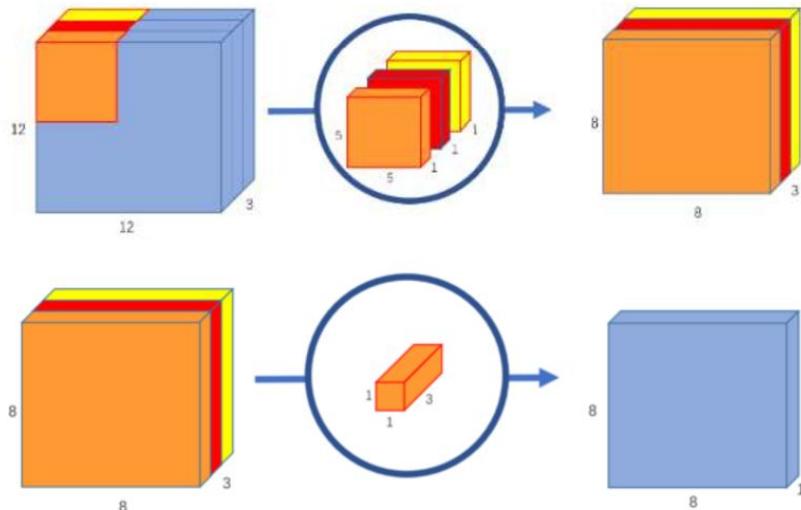
## Simple convolution

# But why?

## Depthwise Convolutional layer:

- Each channel of the input tensor is convolved separately using a  $5 \times 5$  kernel.
- There are 3 input channels and a depth multiplier of 1, so there are 3 depthwise filters in total.

$$3 (5 \times 5 + 1)$$



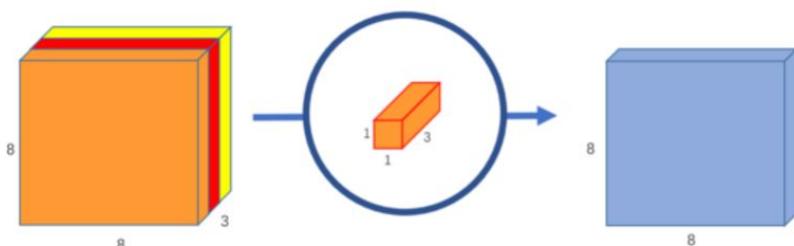
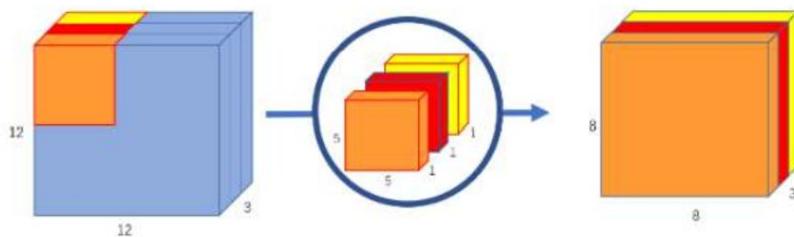
## Pointwise Convolutional layer:

- There are 256 output filters, and the input has 3 channels.
- Each filter has a kernel size of  $1 \times 1 \times 3$  (as 3 input channels). Plus one bias term for each filter.

# But why?

## Depthwise Convolutional layer:

- Each channel of the input tensor is convolved separately using a  $5 \times 5$  kernel.
- There are 3 input channels and a depth multiplier of 1, so there are 3 depthwise filters in total.
- Each depthwise filter has  $5 \times 5 = 25$  parameters. Plus one bias term for each filter.
- Total number of parameters in the Depthwise Convolutional layer is  $3 \times (25+1) = 78$ .



## Pointwise Convolutional layer:

- There are 256 output filters, and the input has 3 channels.
- Each filter has a kernel size of  $1 \times 1 \times 3$  (as 3 input channels). Plus one bias term for each filter.
- Total number of parameters in the Pointwise Convolutional layer is  $256 \times (1 \times 1 \times 3 + 1) = 1024$ .

256  $(1 \times 1 \times 3 + 1)$

# Depthwise Separable Convolution

- **Fewer parameters:** Compared to a standard convolution, depthwise separable convolution requires fewer parameters to learn.
- **Computationally efficient:** The depthwise convolution requires fewer computations than a standard convolution
- **Improved performance:** Depthwise separable convolution can achieve similar or better performance than a standard convolution.

# ImageNet



## Large Scale Visual Recognition Challenge

The Image Classification Challenge:

1,000 object classes

1,431,167 images

- [ILSVRC 2017](#)
- [ILSVRC 2016](#)
- [ILSVRC 2015](#)
- [ILSVRC 2014](#)
- [ILSVRC 2013](#)
- [ILSVRC 2012](#)
- [ILSVRC 2011](#)
- [ILSVRC 2010](#)

# ImageNet

## IMAGENET Large Scale Visual Recognition Challenge

### Imagenet: A large-scale hierarchical image database

[J Deng](#), [W Dong](#), [R Socher](#), [LJ Li](#), [K Li](#)... - 2009 IEEE conference ..., 2009 - ieeexplore.ieee.org

... We introduce here a new **database** called "**ImageNet**", a **large scale** ontology of ... In this paper, we introduce a new **image database** called "**ImageNet**", a **large-scale** ontology of **images**. ...

[☆ Save](#) [✉ Cite](#) [Cited by 51478](#) [Related articles](#) [All 33 versions](#)

# ImageNet

**IMAGENET Large Scale Visual Recognition Challenge**

**Fei-Fei Li**



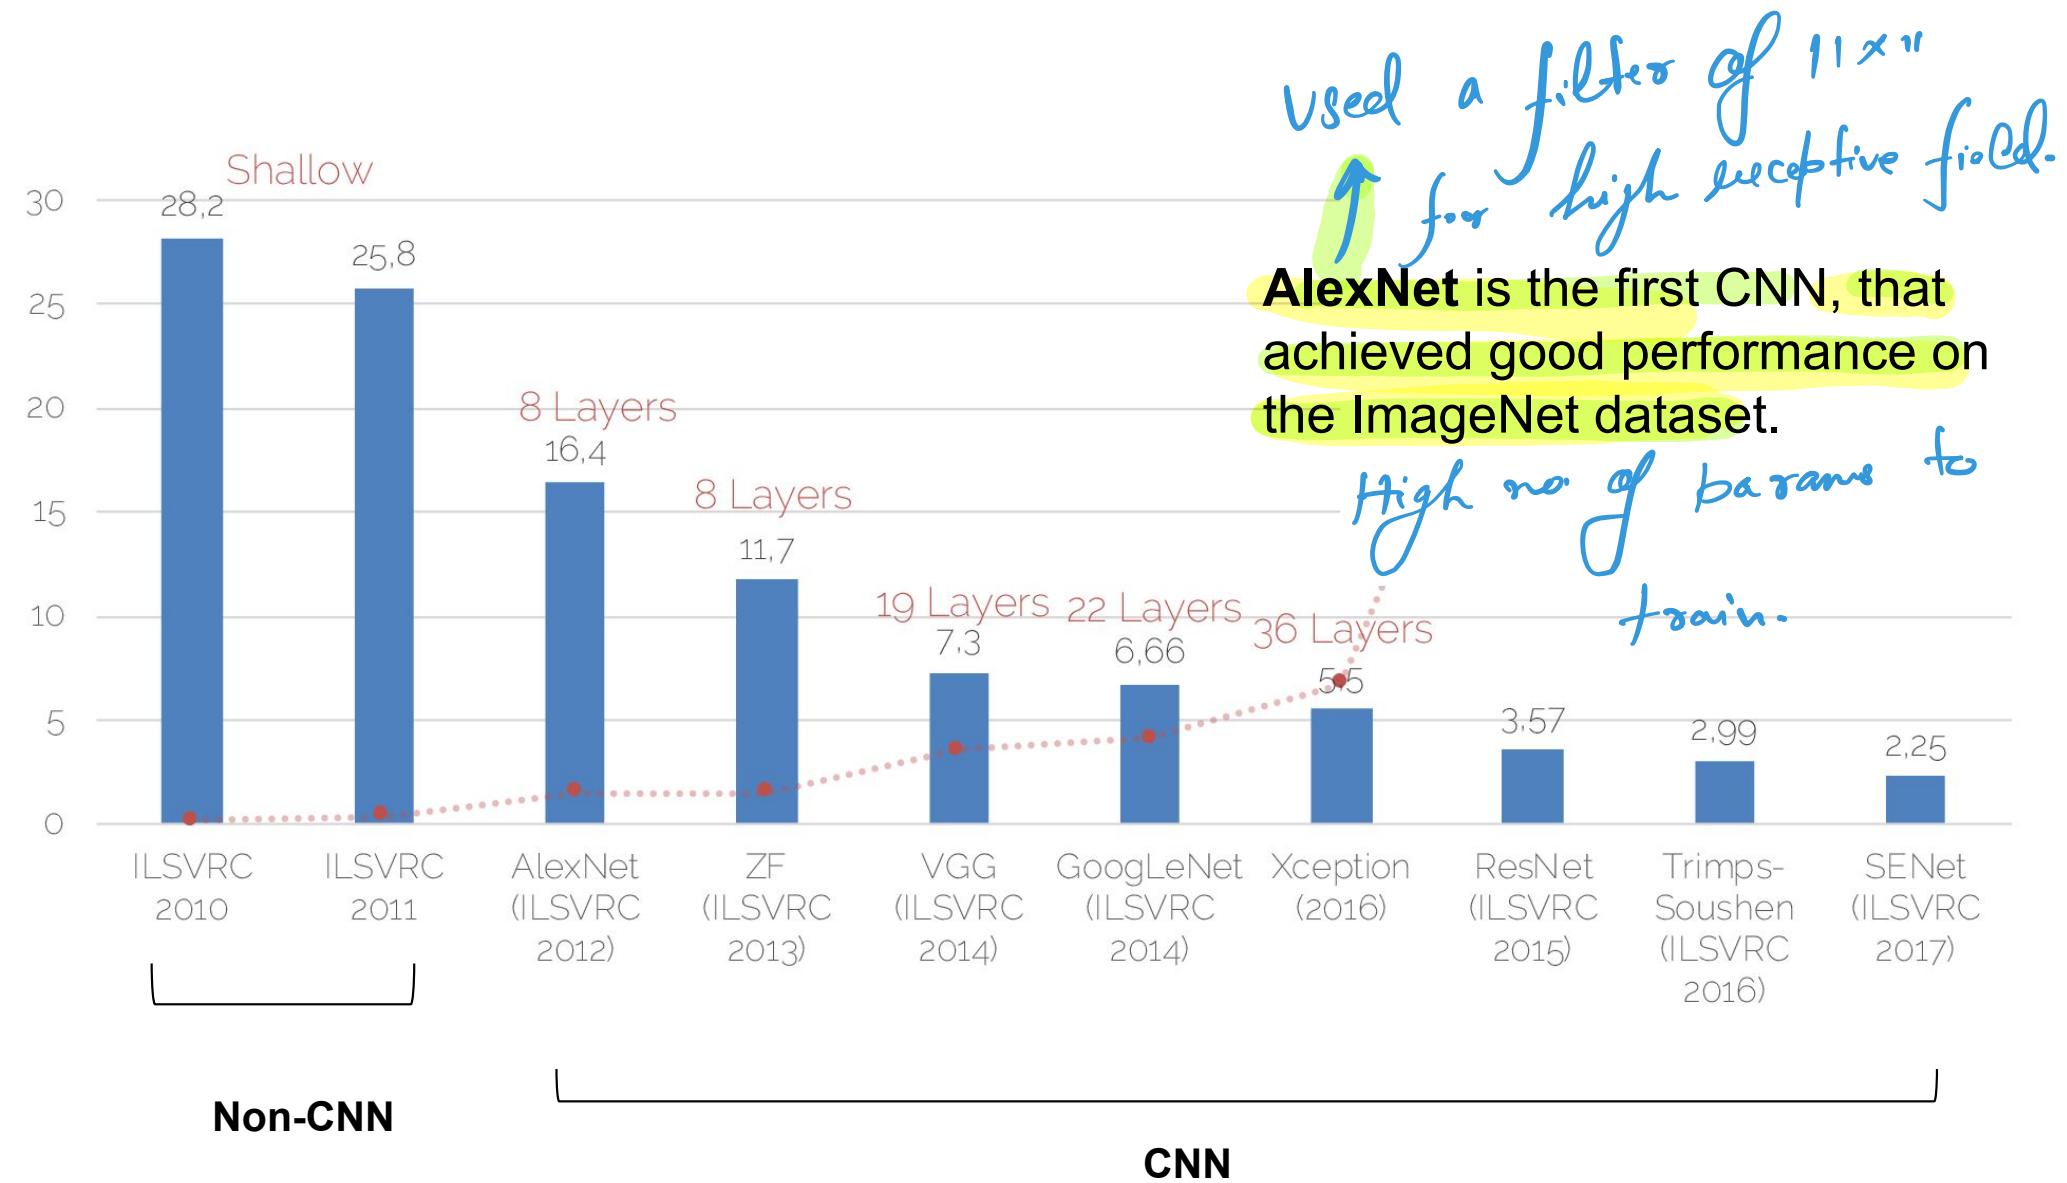
AI researcher **Fei-Fei Li** began working on the idea for ImageNet in 2006.

# Classical Architecture

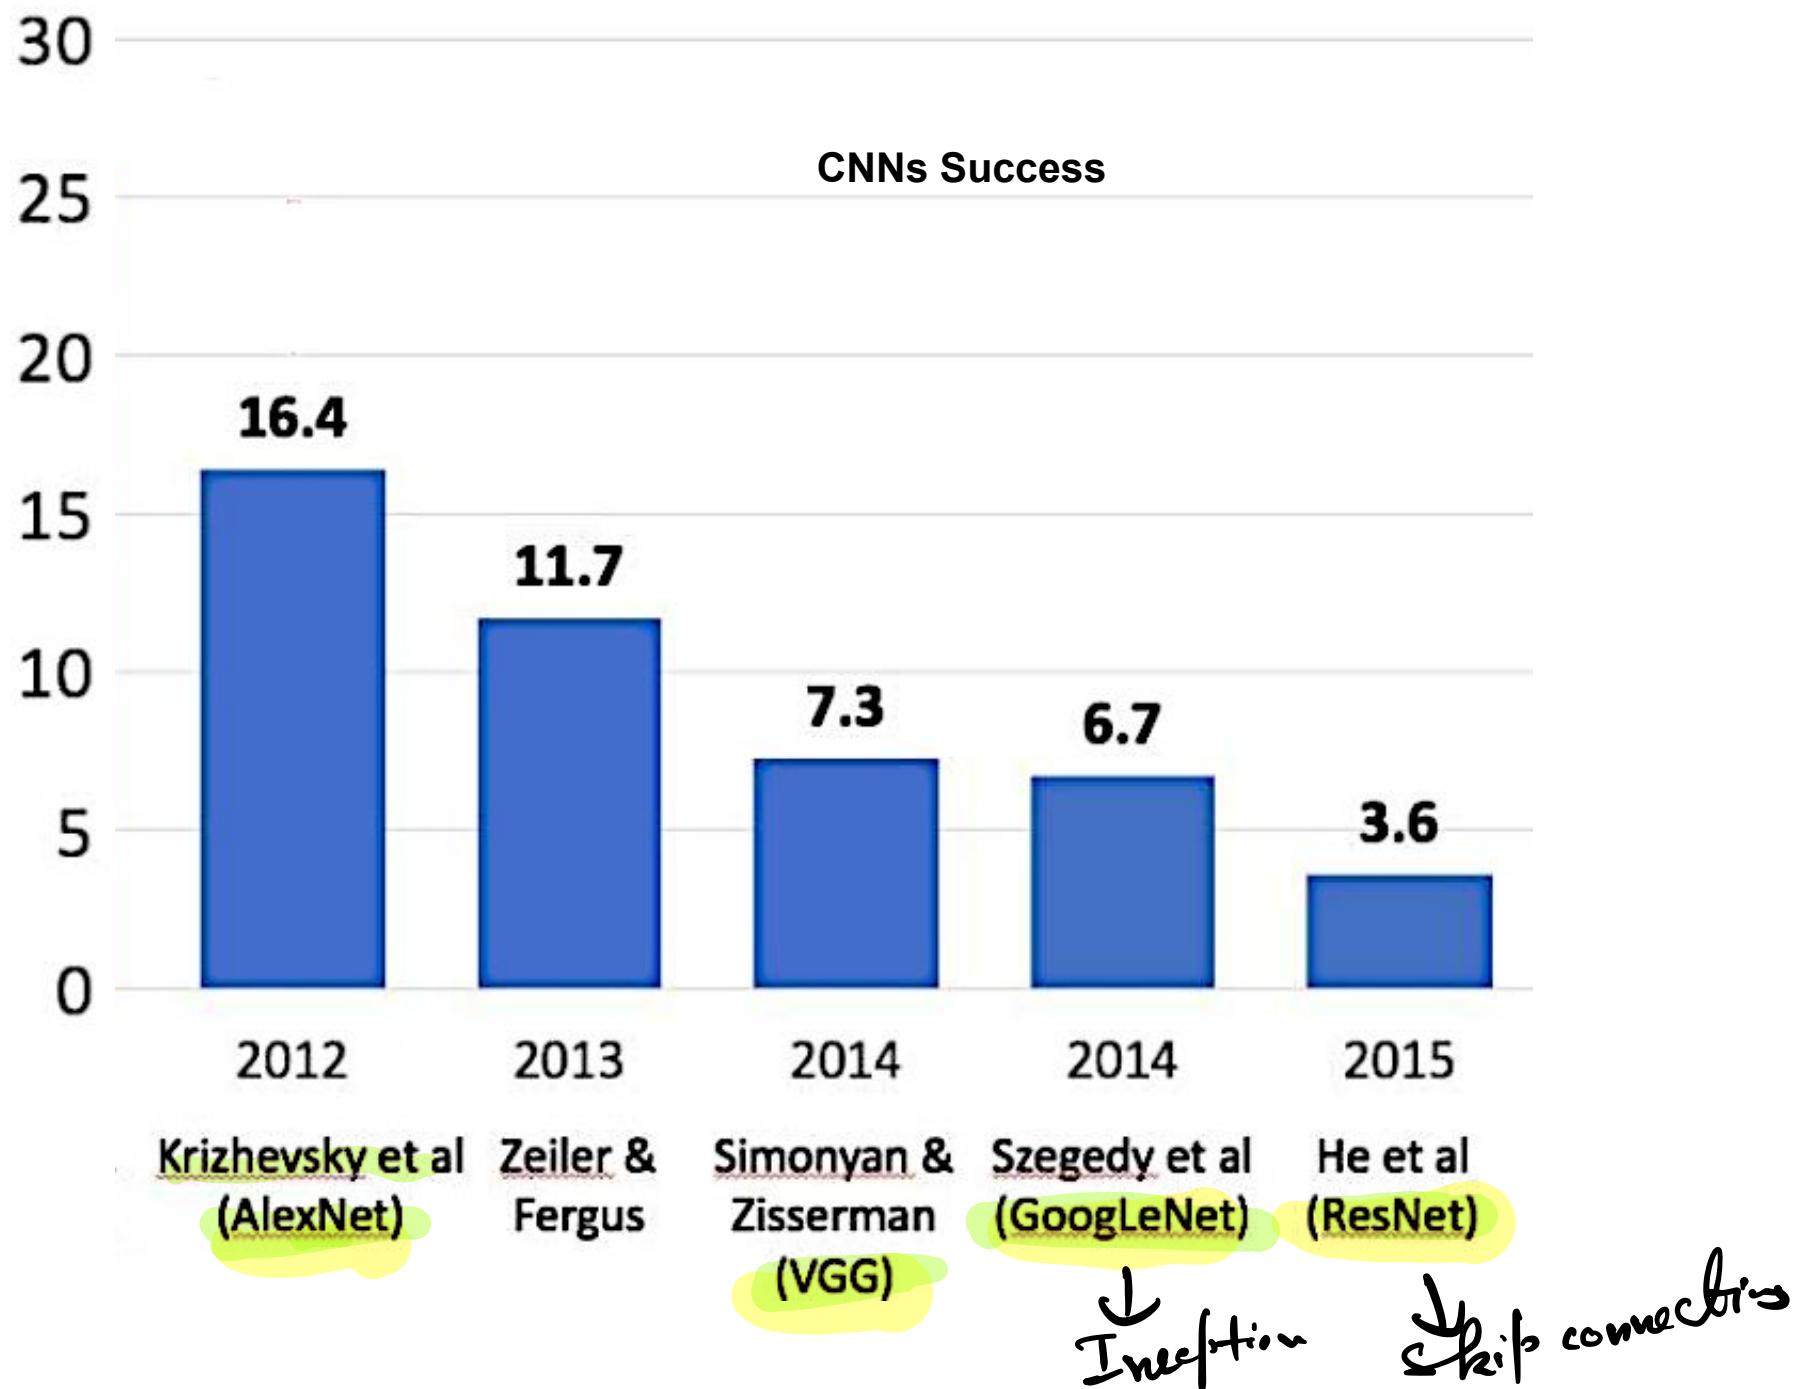
- Alexnet
- VGG Network
- ResNet
- Inception Net

11x11  
 $S=4$

# Revolution of depth (ImageNet Benchmark)



# Alexnet



# Alexnet

- The architecture of AlexNet was inspired by LeNet.
- AlexNet is a convolutional neural network (CNN) architecture designed by Alex Krizhevsky in 2012.

96 11x11



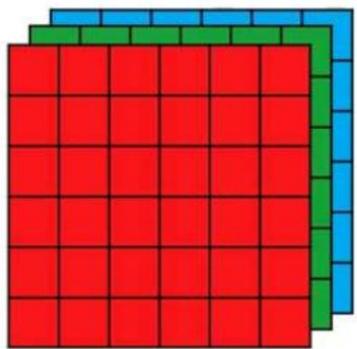
• The architecture of AlexNet was inspired by LeNet.

• AlexNet is a convolutional neural network (CNN) architecture designed by Alex Krizhevsky in 2012.

• AlexNet was one of the first CNN to achieve a good improvement in image classification performance on the ImageNet dataset.

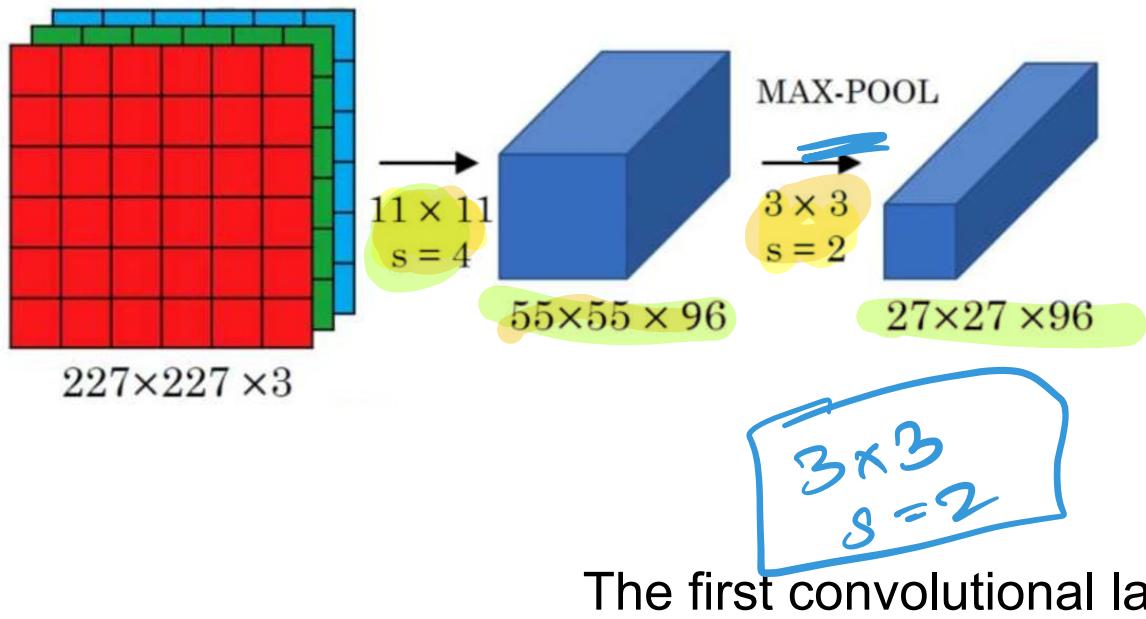
• AlexNet helped to popularize deep learning and convolutional neural networks.

# AlexNet

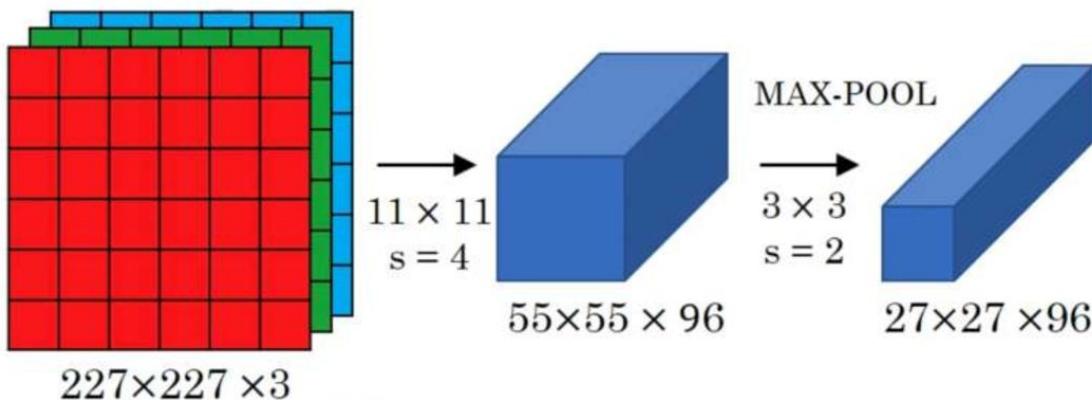


227×227 ×3

# AlexNet



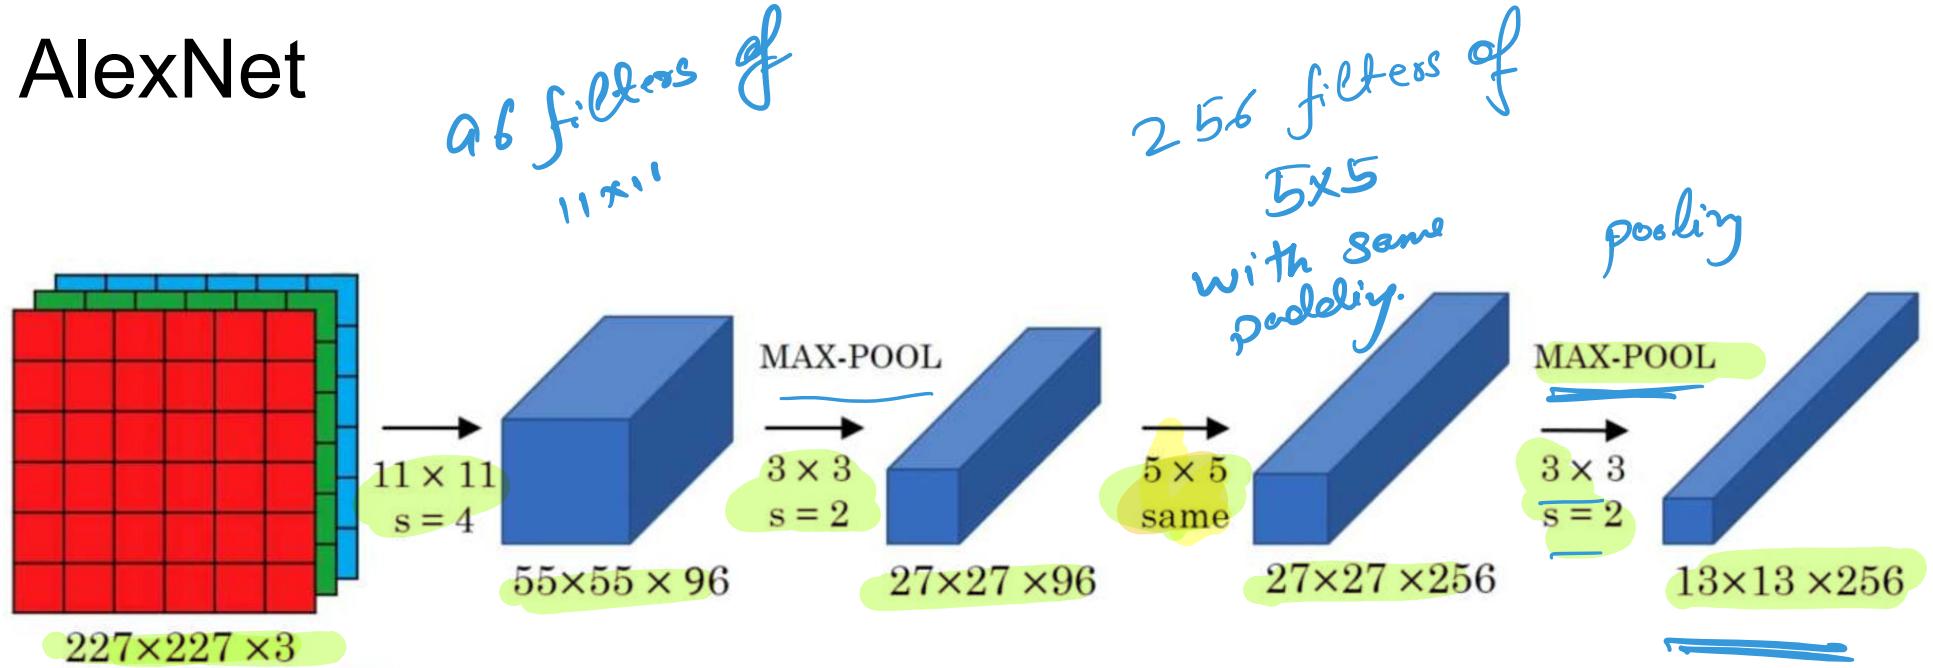
# AlexNet



The first convolutional layer has 96 filters

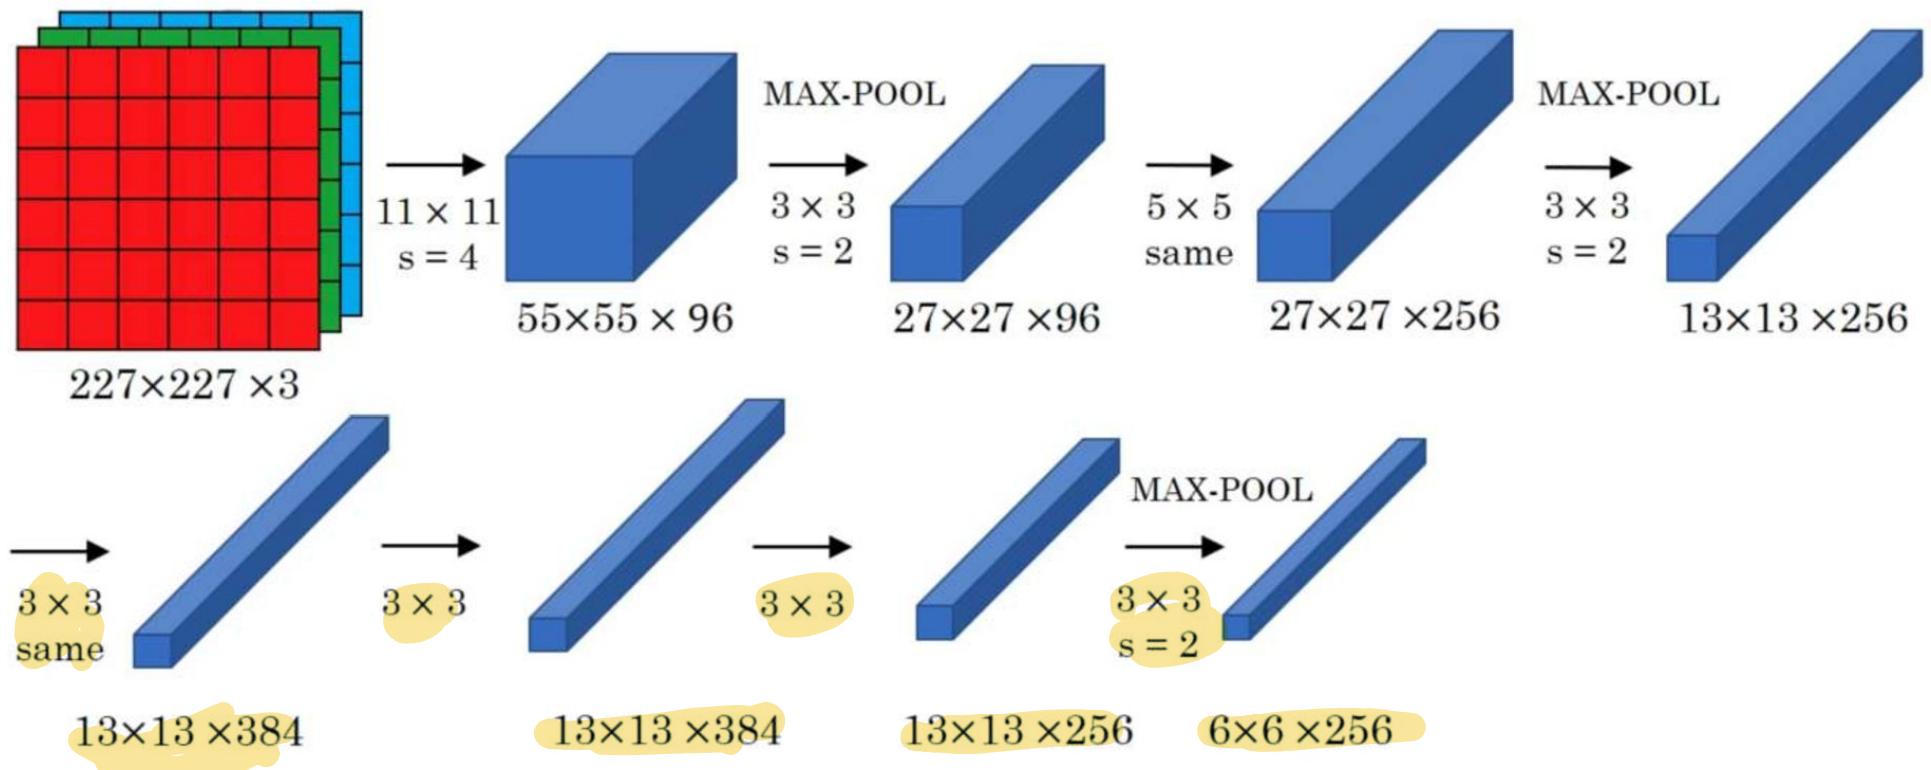
Convolution output is passed through a ReLU activation function and then max pooled with a window size of  $3 \times 3$  and a stride of 2 pixels.

# AlexNet



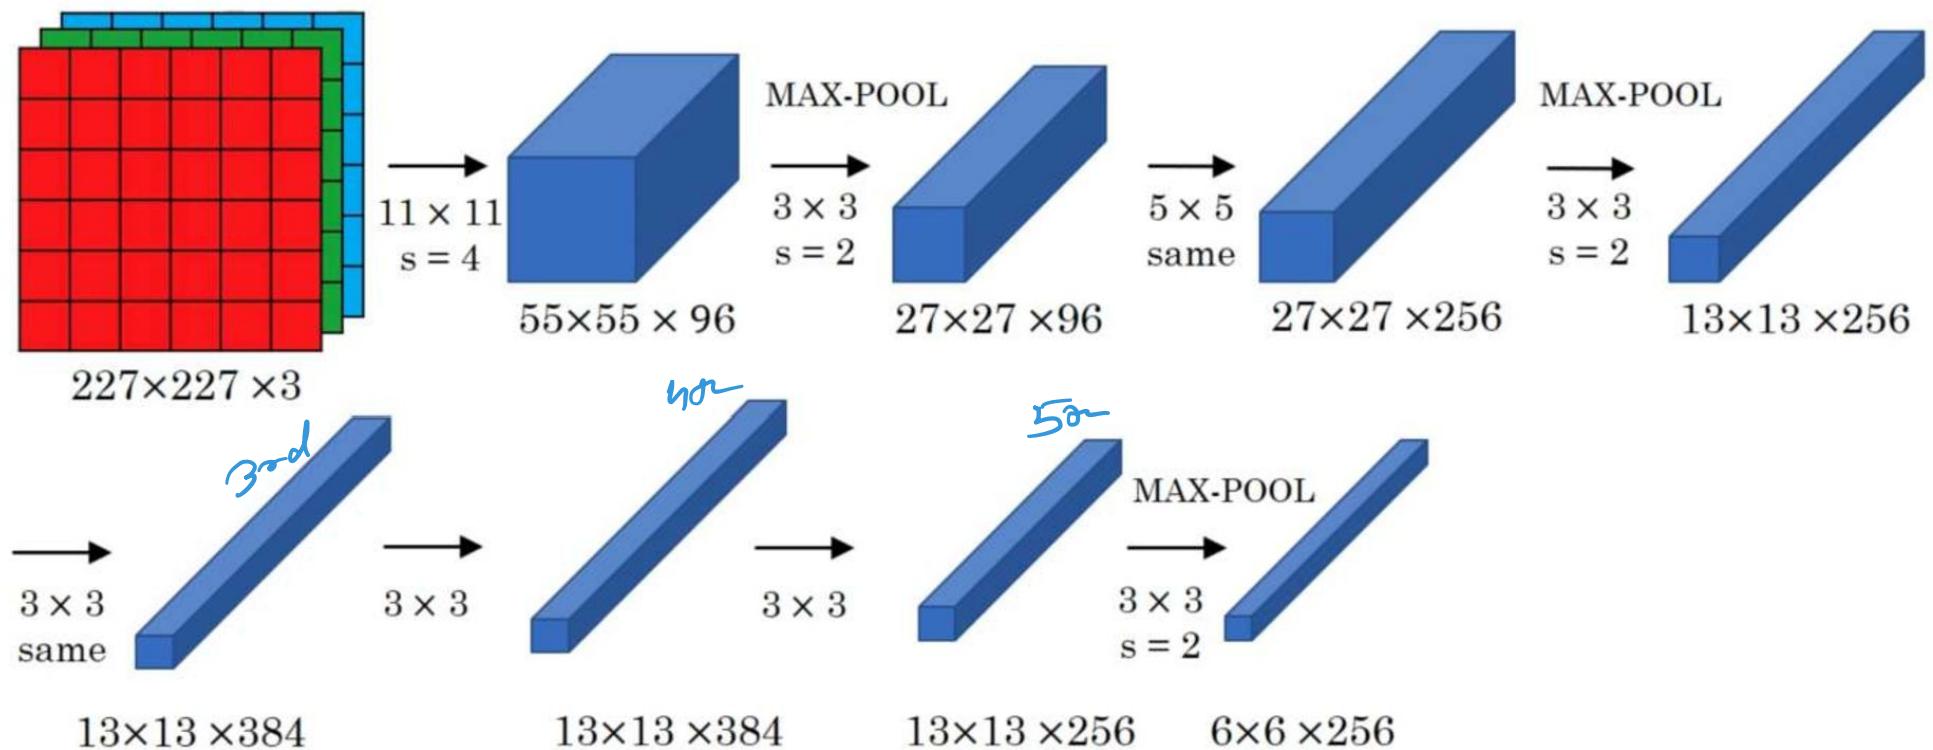
the second convolutional layer has 256 filters.

# AlexNet



Next convolution layers applies 384 filters and 256 filters

# AlexNet

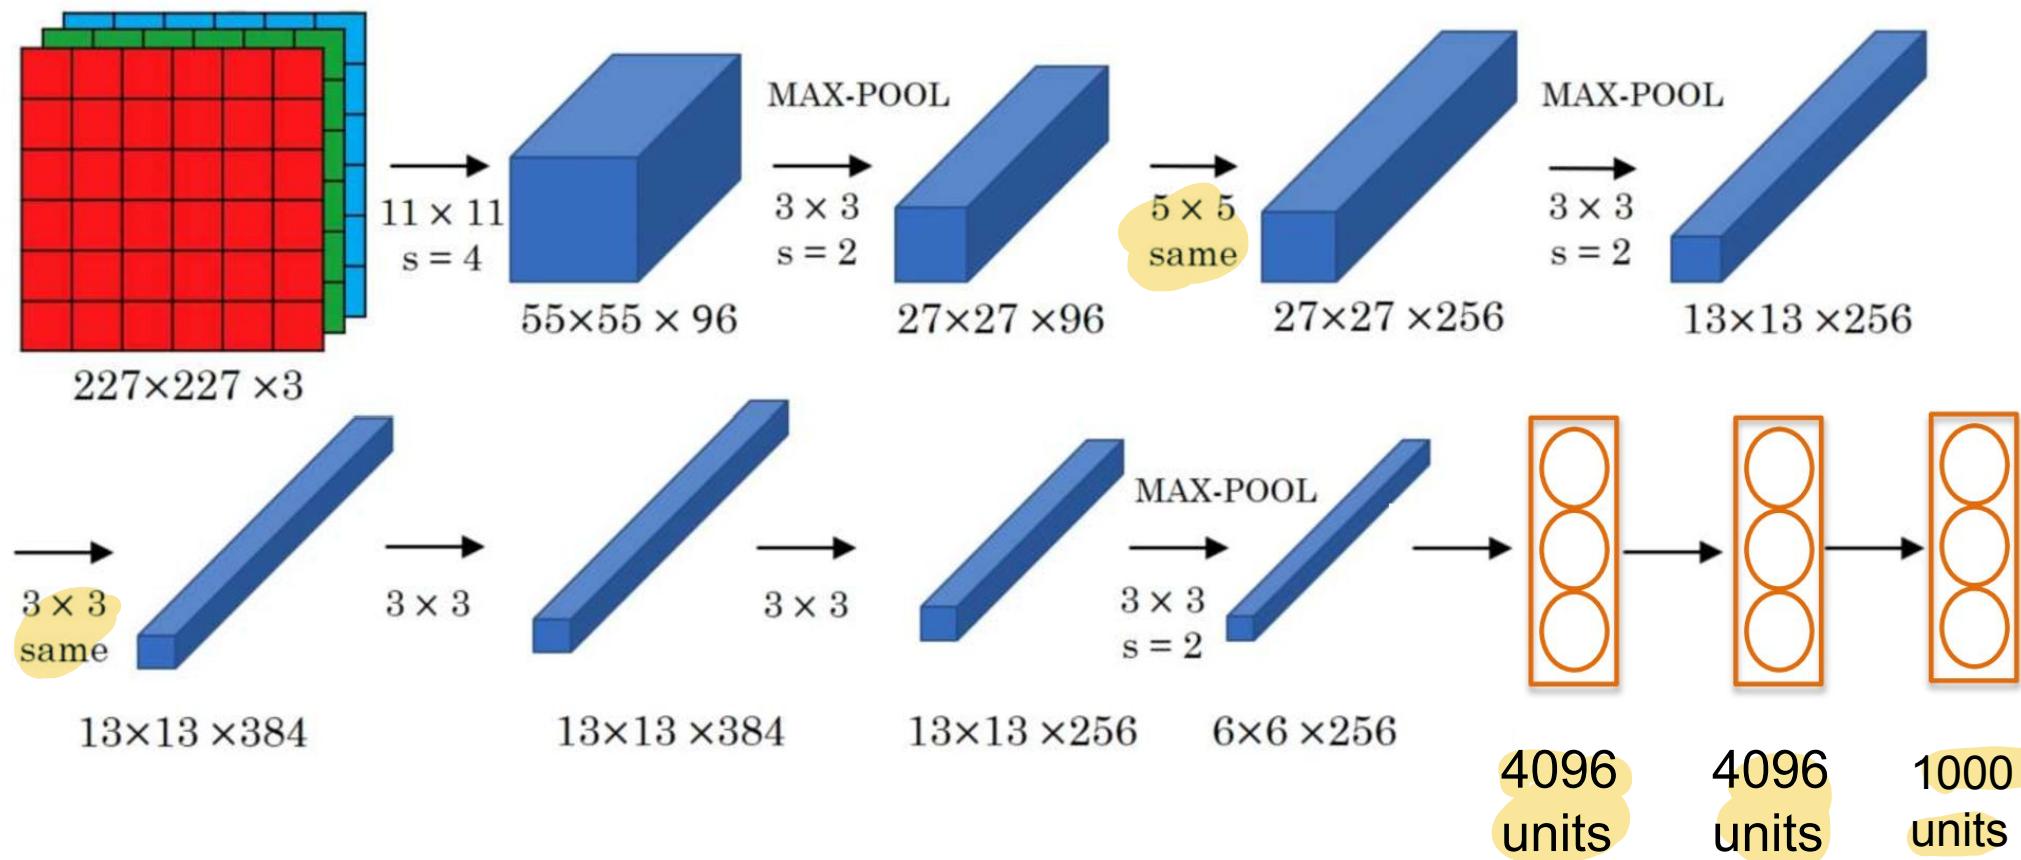


AlexNet uses max pooling after the first, second, and fifth convolutional layers to reduce the spatial dimensions of the feature maps.

1

2

# AlexNet



At last we have Fully Connected layer and softmax layer

# AlexNet

96 11x11

256 5x5

384 3x3

384 3x3

Let Convolutional Layer denoted by **CL**. AlexNet has following layers.

- CL 1: applies 96 filters of size 11x11 to input image with a stride of 4 pixels.
- CL 2: applies 256 filters of size 5x5 to the output of the first max pooling layer. *Same padding* *max pool*
- CL 3: applies 384 filters of size 3x3 to the output of the second max pooling layer.
- CL 4: applies 384 filters of size 3x3 to the output of the third CL.
- CL 5: applies 256 filters of size 3x3 to the output of the fourth CL.

*max pooling.*

Do you observe a pattern in number of filters

Alexnet input 227x227x3

11x11 96 filters s=4

max pooling 3x3 s=2

256 5x5 filters  
max pooling 3x3 s=2

384 3x3 padding s=1  
384 " " "

# AlexNet

256 3x3  
max pooling  
Dense Layer.

Let Convolutional layer denoted by CL. AlexNet has following layers.

- CL 1: applies 96 filters of size 11x11 to input image with a stride of 4 pixels.
- CL 2: applies 256 filters of size 5x5 to the output of the first max pooling layer.
- CL 3: applies 384 filters of size 3x3 to the output of the second max pooling layer.
- CL 4: applies 384 filters of size 3x3 to the output of the third CL.
- CL 5: applies 256 filters of size 3x3 to the output of the fourth CL.

Do you observe a pattern in number of filters

# AlexNet

## Fully connected layers

- FC 1: This layer has 4096 units. The output is passed through a ReLU activation function and then subject to dropout regularization.
- FC 2: This layer is similar to the first fully connected layer, with 4096 units, a ReLU activation function, and dropout regularization.

## Output layer:

- This layer has 1000 units (i.e., number of classes in the ImageNet dataset).
- The output is passed through a softmax activation function to produce the final class probabilities.

# AlexNet

```
from keras.layers import Input, Conv2D, MaxPooling2D, Flatten, Dense, Dropout
from keras.layers.normalization import BatchNormalization
from keras.models import Model

input_shape = (227, 227, 3) # Input shape of the image

# Define the input layer
inputs = Input(shape=input_shape)
```

# AlexNet

```
from keras.layers import Input, Conv2D, MaxPooling2D, Flatten, Dense, Dropout
from keras.layers.normalization import BatchNormalization
from keras.models import Model

input_shape = (227, 227, 3) # Input shape of the image

# Define the input layer
inputs = Input(shape=input_shape)

# First convolutional layer, 96 filters, kernel size of 11x11 and stride of 4x4, followed by ReLU
conv1 = Conv2D(filters=96, kernel_size=(11, 11), strides=(4, 4), activation='relu')(inputs)

# Max pooling layer with pool size of 3x3 and stride of 2x2
pool1 = MaxPooling2D(pool_size=(3, 3), strides=(2, 2))(conv1)

# Batch normalization layer
bn1 = BatchNormalization()(pool1)
```

# AlexNet

```
from keras.layers import Input, Conv2D, MaxPooling2D, Flatten, Dense, Dropout
from keras.layers.normalization import BatchNormalization
from keras.models import Model

input_shape = (227, 227, 3) # Input shape of the image

# Define the input layer
inputs = Input(shape=input_shape)

# First convolutional layer, 96 filters, kernel size of 11x11 and stride of 4x4, followed by ReLU
conv1 = Conv2D(filters=96, kernel_size=(11, 11), strides=(4, 4), activation='relu')(inputs)

# Max pooling layer with pool size of 3x3 and stride of 2x2
pool1 = MaxPooling2D(pool_size=(3, 3), strides=(2, 2))(conv1)

# Batch normalization layer
bn1 = BatchNormalization()(pool1)

# Second convolutional layer with 256 filters, kernel size of 5x5 and padding of same, followed by
# ReLU activation
conv2 = Conv2D(filters=256, kernel_size=(5, 5), padding='same', activation='relu')(bn1)

# Max pooling layer with pool size of 3x3 and stride of 2x2
pool2 = MaxPooling2D(pool_size=(3, 3), strides=(2, 2))(conv2)

# Batch normalization layer
bn2 = BatchNormalization()(pool2)
```

# AlexNet

```
# Third convolutional layer with 384 filters, kernel size of 3x3 and padding of same, followed by ReLU
conv3 = Conv2D(filters=384, kernel_size=(3, 3), padding='same', activation='relu')(bn2)
# Fourth convolutional layer with 384 filters, kernel size of 3x3 and padding of same, followed by ReLU
conv4 = Conv2D(filters=384, kernel_size=(3, 3), padding='same', activation='relu')(conv3)
```

# AlexNet

```
# Third convolutional layer with 384 filters, kernel size of 3x3 and padding of same, followed by ReLU
conv3 = Conv2D(filters=384, kernel_size=(3, 3), padding='same', activation='relu')(bn2)
# Fourth convolutional layer with 384 filters, kernel size of 3x3 and padding of same, followed by ReLU
conv4 = Conv2D(filters=384, kernel_size=(3, 3), padding='same', activation='relu')(conv3)

# Fifth convolutional layer, 256 filters, kernel size of 3x3 and padding of same, followed by ReLU
conv5 = Conv2D(filters=256, kernel_size=(3, 3), padding='same', activation='relu')(conv4)
# Max pooling layer with pool size of 3x3 and stride of 2x2
pool5 = MaxPooling2D(pool_size=(3, 3), strides=(2, 2))(conv5)
# Batch normalization layer
bn3 = BatchNormalization()(pool5)
```

# AlexNet

```
# Third convolutional layer with 384 filters, kernel size of 3x3 and padding of same, followed by ReLU
conv3 = Conv2D(filters=384, kernel_size=(3, 3), padding='same', activation='relu')(bn2)
# Fourth convolutional layer with 384 filters, kernel size of 3x3 and padding of same, followed by ReLU
conv4 = Conv2D(filters=384, kernel_size=(3, 3), padding='same', activation='relu')(conv3)

# Fifth convolutional layer, 256 filters, kernel size of 3x3 and padding of same, followed by ReLU
conv5 = Conv2D(filters=256, kernel_size=(3, 3), padding='same', activation='relu')(conv4)
# Max pooling layer with pool size of 3x3 and stride of 2x2
pool5 = MaxPooling2D(pool_size=(3, 3), strides=(2, 2))(conv5)
# Batch normalization layer
bn3 = BatchNormalization()(pool5)

# Flatten layer
flatten = Flatten()(bn3)

# First fully connected layer with 4096 units, followed by ReLU activation and dropout
fc1 = Dense(units=4096, activation='relu')(flatten)
dropout1 = Dropout(0.5)(fc1)
# Second fully connected layer with 4096 units, followed by ReLU activation and dropout
fc2 = Dense(units=4096, activation='relu')(dropout1)
dropout2 = Dropout(0.5)(fc2)

# Output layer with 1000 units and softmax activation
outputs = Dense(units=1000, activation='softmax')(dropout2)
```

# LeNet and AlexNet

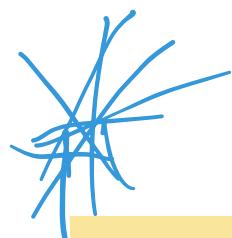
- AlexNet is deeper and more complex than LeNet, with more layers and more parameters.

# LeNet and AlexNet

- AlexNet is deeper and more complex than LeNet, with more layers and more parameters.
- AlexNet has larger filters in initial layers (e.g., 11x11 in the first layer),  
while LeNet uses smaller filters (e.g., 5x5).

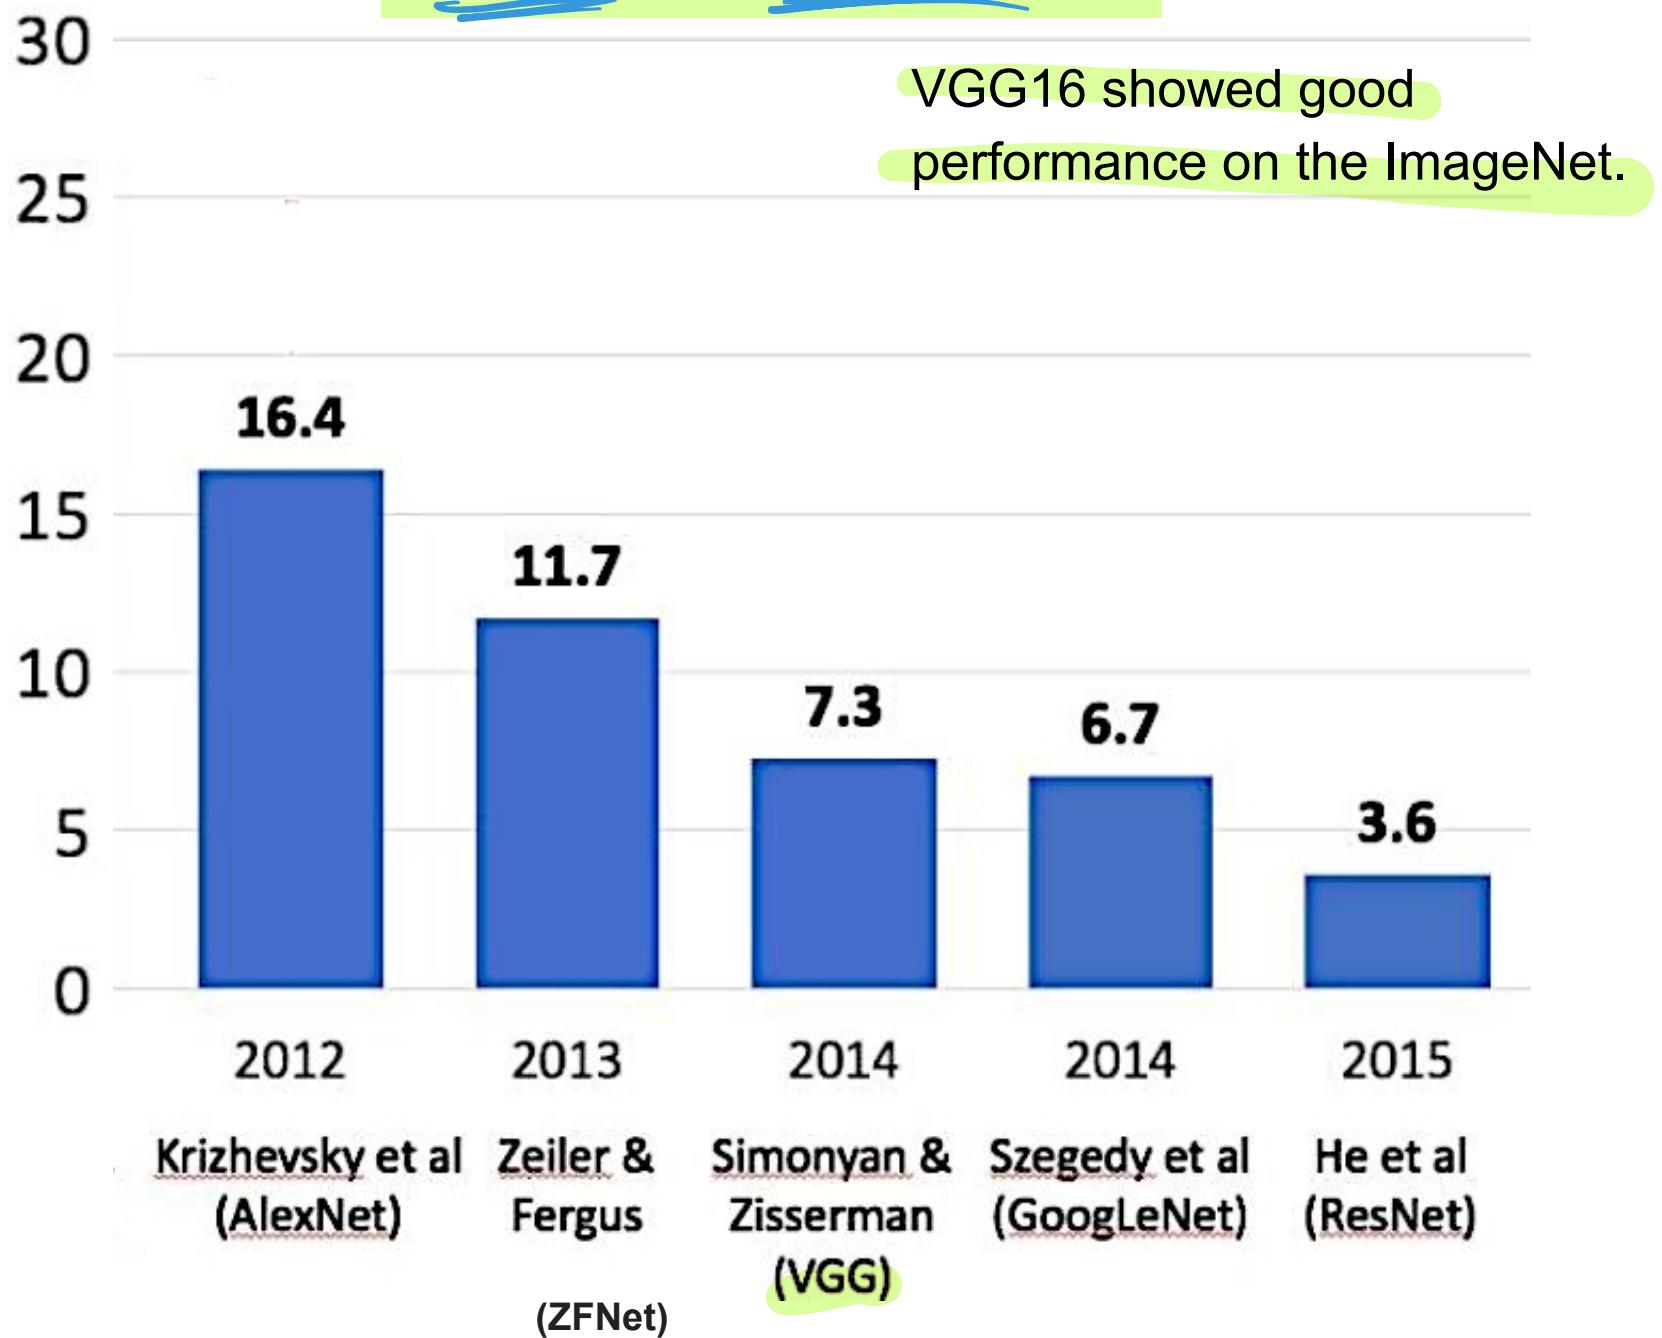
# LeNet and AlexNet

- AlexNet is deeper and more complex than LeNet, with more layers and more parameters.
- AlexNet has larger filters in initial layers (e.g.,  $11 \times 11$  in the first layer), while LeNet uses smaller filters (e.g.,  $5 \times 5$ ).
- AlexNet achieved state-of-the-art performance on the ImageNet dataset, while LeNet was designed and tested on a smaller handwritten digit recognition task.



# VGG Network

# VGG Network



# VGG Networks

- VGG stands for Visual Geometry Group (University of Oxford) that developed a family of CNN architectures for image classification.

# VGG Networks

- VGG stands for Visual Geometry Group (University of Oxford) that developed a family of CNN architectures for image classification.
- The original VGG model, VGG16, has 16 layers including 13 convolutional layers and 3 fully connected layers.

3 fc layers are there in  
every VGG.

# VGG Networks

- VGG stands for Visual Geometry Group (University of Oxford) that developed a family of CNN architectures for image classification.
- The original VGG model, VGG16, has 16 layers including 13 convolutional layers and 3 fully connected layers.
- VGG19 is with 19 layers including 16 convolutional layers and 3 fully connected layers.

# VGG Networks

- VGG stands for Visual Geometry Group (University of Oxford) that developed a family of CNN architectures for image classification.
- The original VGG model, VGG16, has 16 layers including 13 convolutional layers and 3 fully connected layers.
- VGG19 is with 19 layers including 16 convolutional layers and 3 fully connected layers.
- VGG family also includes VGG11 and VGG13, with fewer convolutional layers than VGG16 and VGG19.

# VGG Networks

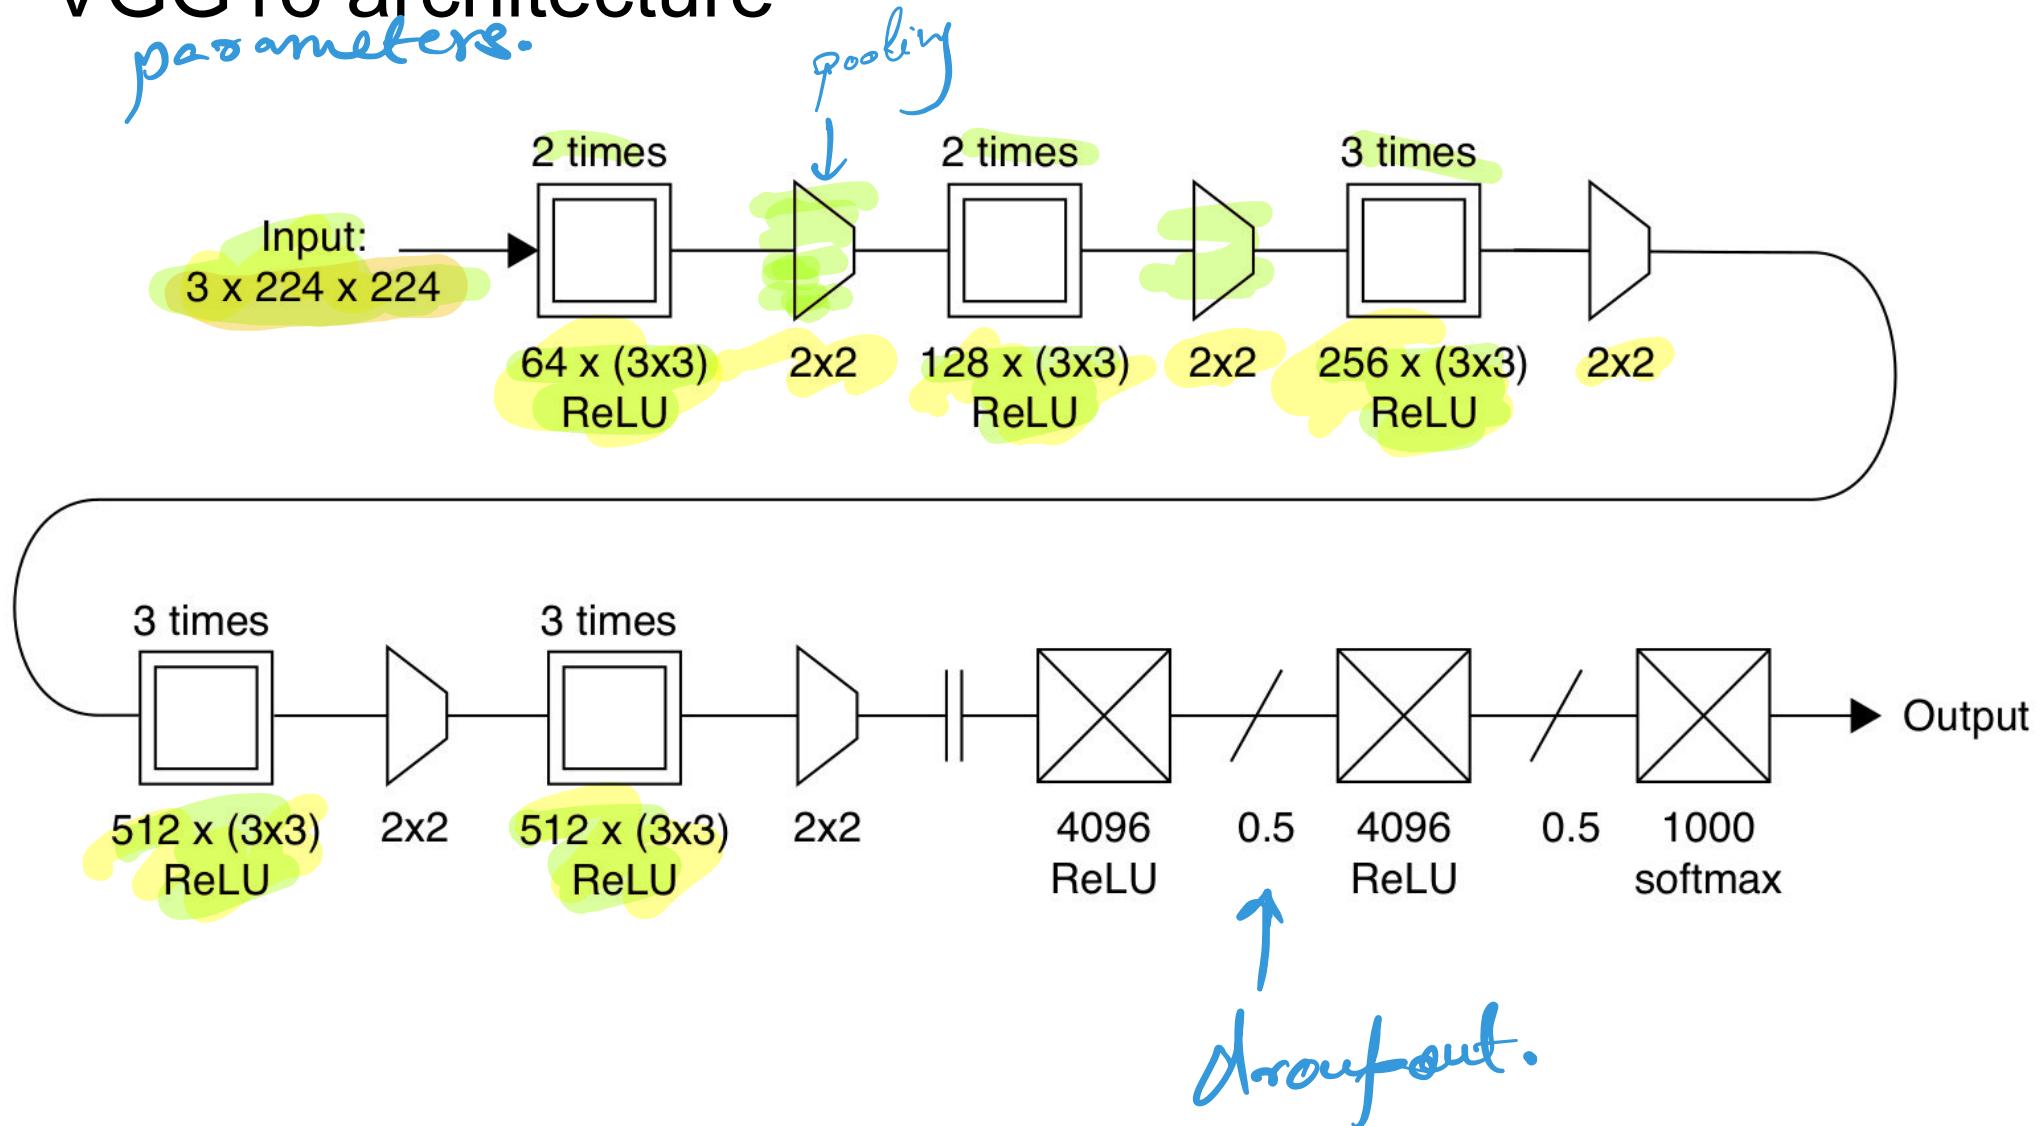
- VGG stands for Visual Geometry Group (University of Oxford) that developed a family of CNN architectures for image classification.
- The original VGG model, VGG16, has 16 layers including 13 convolutional layers and 3 fully connected layers.
- VGG19 is with 19 layers including 16 convolutional layers and 3 fully connected layers.
- VGG family also includes VGG11 and VGG13, with fewer convolutional layers than VGG16 and VGG19.
- VGG family generalize well to many tasks such as classification, object detection, segmentation, style transfer, and transfer learning.

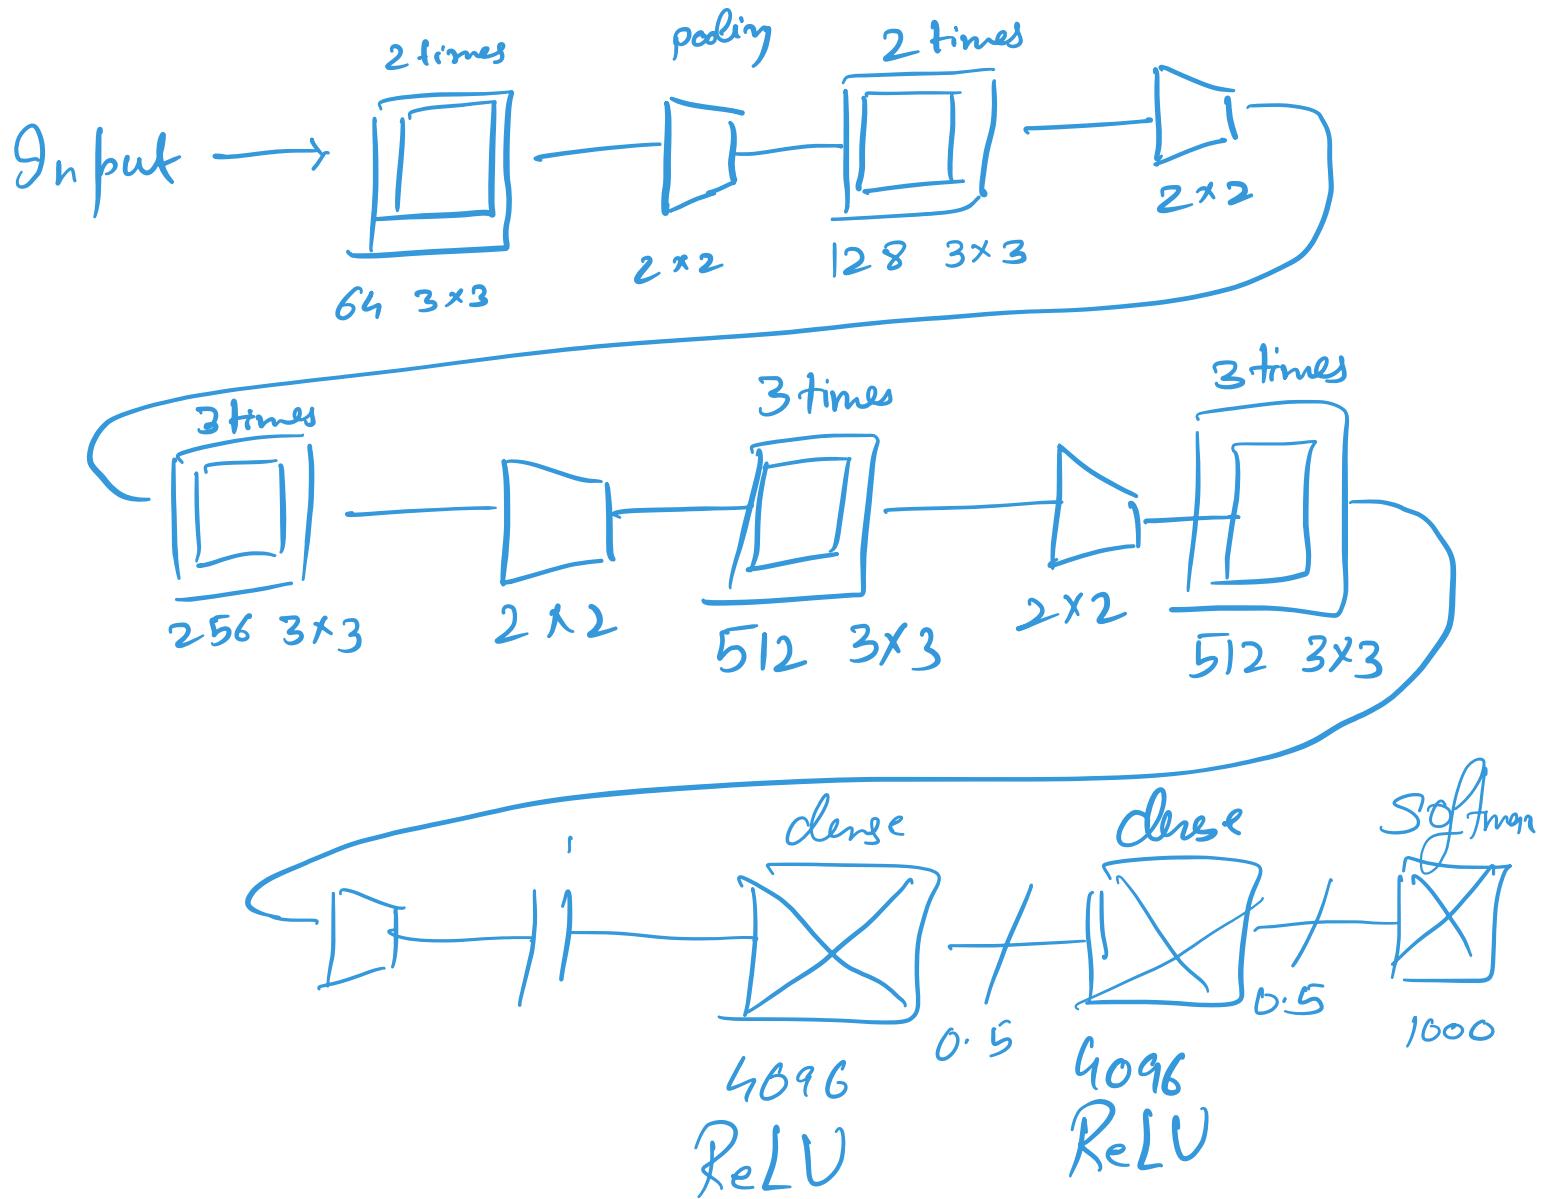
what VGG does better than Alexnet  
→ Use of  $3 \times 3$  filters multiple times gave

higher receptive field with the use of less parameters.

## VGG16 architecture

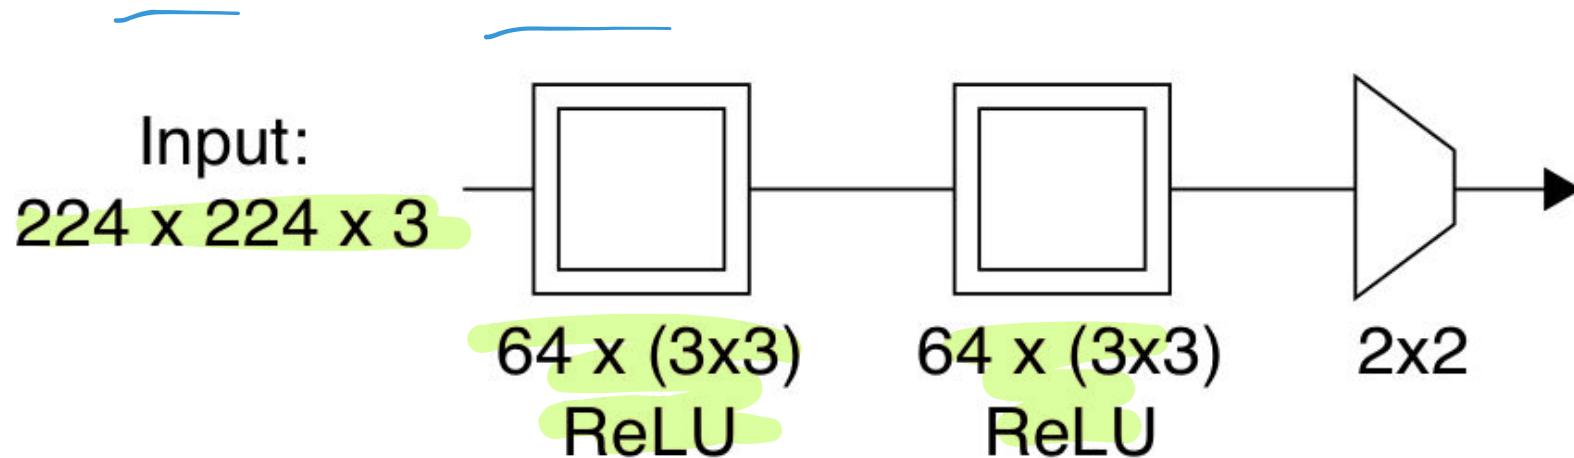
*parameters.*





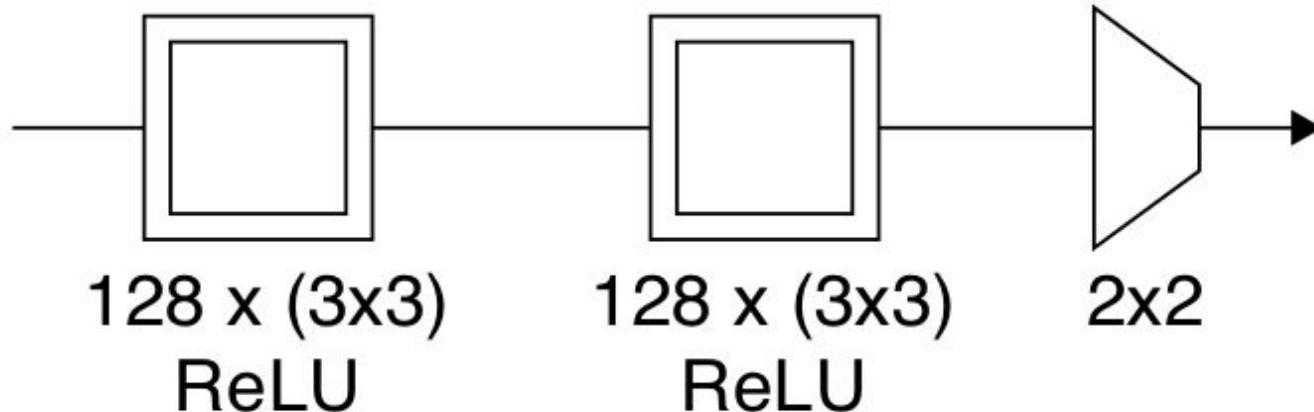
| Dy.

## VGG16 architecture



Group 1 of VGG16. We convolve the input tensor with 64 filters each of size 3 by 3. Then we convolve again with 64 new filters. Finally, we use max pooling to reduce the output tensor's height and width by half.

# VGG16 architecture



*Group 2 of VGG16 is just like the first block  
except that we use 128 filters in each convolution layer rather than 64.*

2  
64

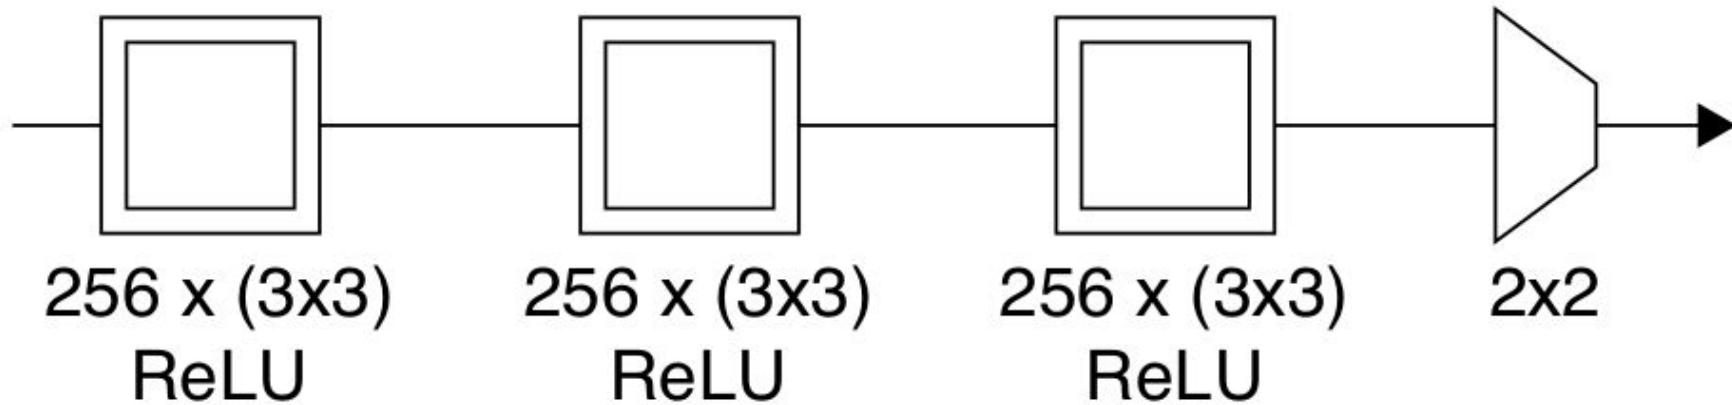
2  
128

3  
256

↓

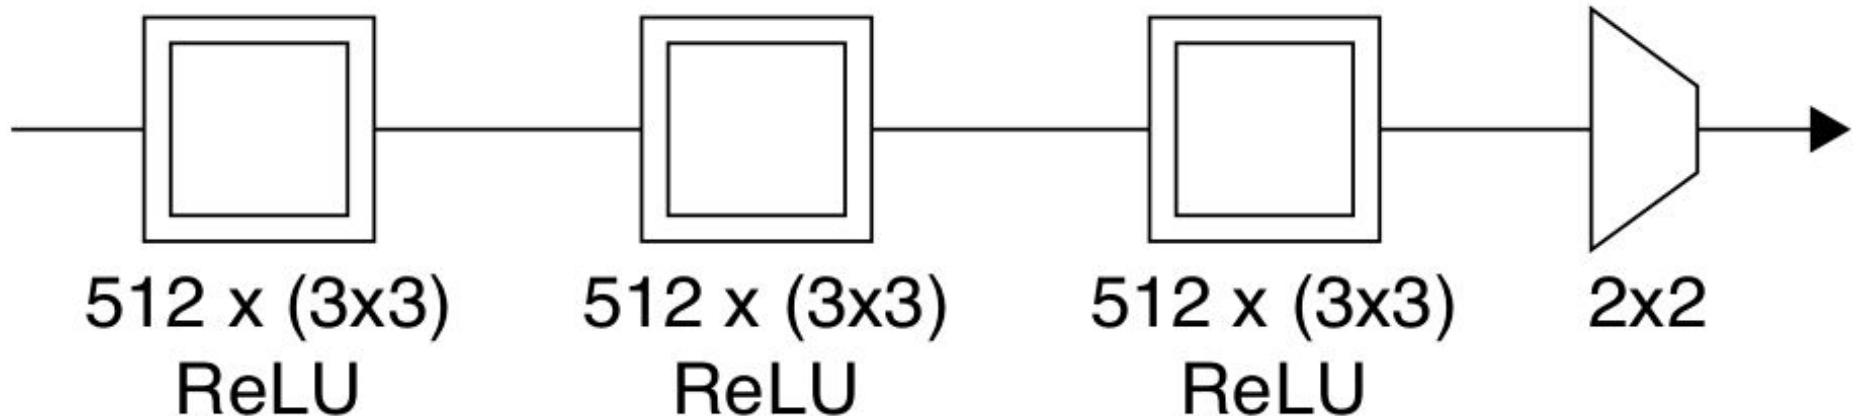
## VGG16 architecture

5 12 3  
↓  
5 52 3



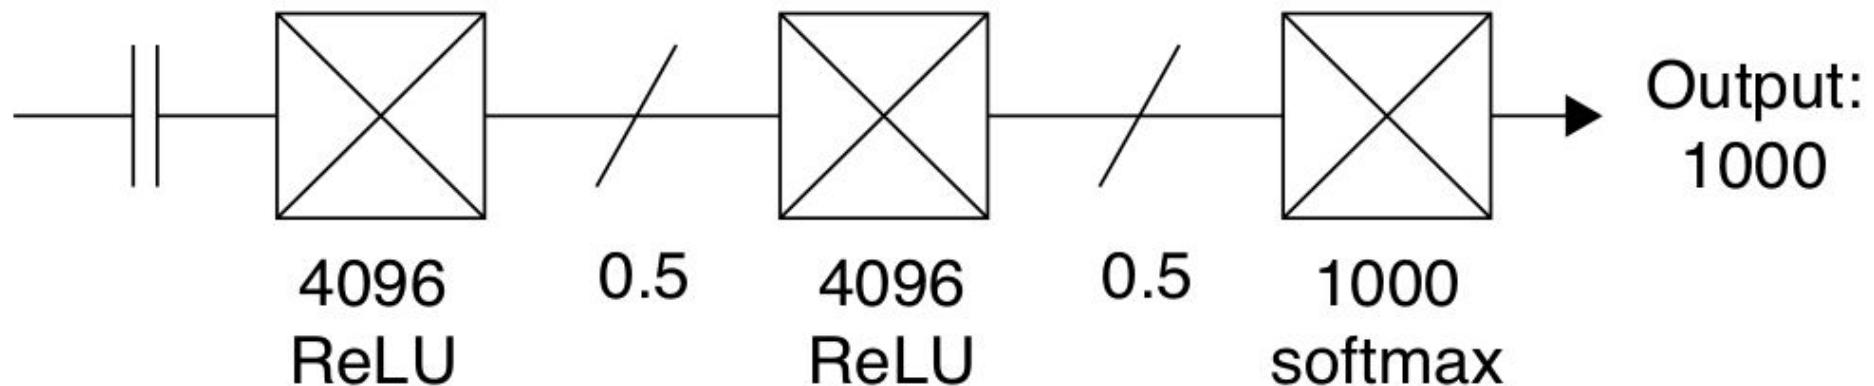
*Group 3 of VGG16 doubles the number of filters again to 256 and repeats the convolution step three times rather than two as before.*

# VGG16 architecture



Groups 4 and 5 of VGG16 are the same. They each have three convolution layers, followed by a two by two max pooling layer.

# VGG16 architecture



*The final steps of processing in VGG16. We flatten the image, then run it through two dense layers each using ReLU, followed by dropout, then through a dense layer with softmax.*

# VGG16 (summary)

- VGG16 has 13 convolutional layers and 3 fully connected layers (16 total).

## VGG16 (summary)

- VGG16 has 13 convolutional layers and 3 fully connected layers (16 total).
- *Convolutional layers* use 3x3 filters, stride 1, same padding, followed by a ReLU (deeper networks & less parameters)

# VGG16 (summary)

- VGG16 has 13 convolutional layers and 3 fully connected layers (16 total).
- *Convolutional layers* use 3x3 filters, stride 1, same padding, followed by a ReLU (deeper networks & less parameters)
- Pooling layers use 2x2 filters with a stride of 2 pixels.

# VGG16 (summary)

- VGG16 has 13 convolutional layers and 3 fully connected layers (16 total).
- *Convolutional layers* use 3x3 filters, stride 1, same padding, followed by a ReLU (deeper networks & less parameters)
- Pooling layers use 2x2 filters with a stride of 2 pixels.
- The first two convolutional layers have 64 filters each, while the remaining layers have 128, 256, 512, and 512 filters, respectively.

# VGG16 (summary)

- VGG16 has 13 convolutional layers and 3 fully connected layers (16 total).
- *Convolutional layers* use 3x3 filters, stride 1, same padding, followed by a ReLU (deeper networks & less parameters)
- Pooling layers use 2x2 filters with a stride of 2 pixels.
- The first two convolutional layers have 64 filters each, while the remaining layers have 128, 256, 512, and 512 filters, respectively.
- Fully connected layers have 4096 units that use ReLU, and output layer has 1000 units corresponding to number of classes in the ImageNet.

# VGG16

```
# Import necessary libraries
import tensorflow as tf
from tensorflow.keras.applications.vgg16 import VGG16
from tensorflow.keras.preprocessing import image
from tensorflow.keras.applications.vgg16 import preprocess_input, decode_predictions
```

# VGG16

```
# Import necessary libraries
import tensorflow as tf
from tensorflow.keras.applications.vgg16 import VGG16
from tensorflow.keras.preprocessing import image
from tensorflow.keras.applications.vgg16 import preprocess_input, decode_predictions

# Load the VGG16 model
model = VGG16(weights='imagenet')
```

# VGG16

```
# Import necessary libraries
import tensorflow as tf
from tensorflow.keras.applications.vgg16 import VGG16
from tensorflow.keras.preprocessing import image
from tensorflow.keras.applications.vgg16 import preprocess_input, decode_predictions

# Load the VGG16 model
model = VGG16(weights='imagenet')

# Load the image you want to classify
img_path = 'tiger_shark.jpeg'
img = image.load_img(img_path, target_size=(224, 224))
```

# VGG16

```
# Import necessary libraries
import tensorflow as tf
from tensorflow.keras.applications.vgg16 import VGG16
from tensorflow.keras.preprocessing import image
from tensorflow.keras.applications.vgg16 import preprocess_input, decode_predictions

# Load the VGG16 model
model = VGG16(weights='imagenet')

# Load the image you want to classify
img_path = 'tiger_shark.jpeg'
img = image.load_img(img_path, target_size=(224, 224))

# Convert the image to an array
x = image.img_to_array(img)
x = tf.expand_dims(x, axis=0)
x = preprocess_input(x)
```

# VGG16

```
# Import necessary libraries
import tensorflow as tf
from tensorflow.keras.applications.vgg16 import VGG16
from tensorflow.keras.preprocessing import image
from tensorflow.keras.applications.vgg16 import preprocess_input, decode_predictions

# Load the VGG16 model
model = VGG16(weights='imagenet')

# Load the image you want to classify
img_path = 'tiger_shark.jpeg'
img = image.load_img(img_path, target_size=(224, 224))

# Convert the image to an array
x = image.img_to_array(img)
x = tf.expand_dims(x, axis=0)
x = preprocess_input(x)

# Use the model to predict the class of the image
preds = model.predict(x)

# Print the top 5 predictions
print('Predicted:', decode_predictions(preds, top=5)[0])
```

# VGG16 Implementation

```
# Import necessary libraries
from tensorflow.keras.models import Model
from tensorflow.keras.layers import Input, Conv2D, MaxPooling2D, Flatten, Dense

def vgg16(input_shape=(224, 224, 3), num_classes=1000):
    input_tensor = Input(shape=input_shape)
```

# VGG16 Implementation

```
# Import necessary libraries
from tensorflow.keras.models import Model
from tensorflow.keras.layers import Input, Conv2D, MaxPooling2D, Flatten, Dense

def vgg16(input_shape=(224, 224, 3), num_classes=1000):
    input_tensor = Input(shape=input_shape)

    # Block 1
    x = Conv2D(64, (3, 3), activation='relu', padding='same', name='block1_conv1')(input_tensor)
    x = Conv2D(64, (3, 3), activation='relu', padding='same', name='block1_conv2')(x)
    x = MaxPooling2D((2, 2), strides=(2, 2), name='block1_pool')(x)
```

# VGG16 Implementation

```
# Import necessary libraries
from tensorflow.keras.models import Model
from tensorflow.keras.layers import Input, Conv2D, MaxPooling2D, Flatten, Dense

def vgg16(input_shape=(224, 224, 3), num_classes=1000):
    input_tensor = Input(shape=input_shape)

    # Block 1
    x = Conv2D(64, (3, 3), activation='relu', padding='same', name='block1_conv1')(input_tensor)
    x = Conv2D(64, (3, 3), activation='relu', padding='same', name='block1_conv2')(x)
    x = MaxPooling2D((2, 2), strides=(2, 2), name='block1_pool')(x)

    # Block 2
    x = Conv2D(128, (3, 3), activation='relu', padding='same', name='block2_conv1')(x)
    x = Conv2D(128, (3, 3), activation='relu', padding='same', name='block2_conv2')(x)
    x = MaxPooling2D((2, 2), strides=(2, 2), name='block2_pool')(x)

    # Block 3
    x = Conv2D(256, (3, 3), activation='relu', padding='same', name='block3_conv1')(x)
    x = Conv2D(256, (3, 3), activation='relu', padding='same', name='block3_conv2')(x)
    x = Conv2D(256, (3, 3), activation='relu', padding='same', name='block3_conv3')(x)
    x = MaxPooling2D((2, 2), strides=(2, 2), name='block3_pool')(x)
```

# VGG16 Implementation

```
# Block 4
x = Conv2D(512, (3, 3), activation='relu', padding='same', name='block4_conv1')(x)
x = Conv2D(512, (3, 3), activation='relu', padding='same', name='block4_conv2')(x)
x = Conv2D(512, (3, 3), activation='relu', padding='same', name='block4_conv3')(x)
x = MaxPooling2D((2, 2), strides=(2, 2), name='block4_pool')(x)

# Block 5
x = Conv2D(512, (3, 3), activation='relu', padding='same', name='block5_conv1')(x)
x = Conv2D(512, (3, 3), activation='relu', padding='same', name='block5_conv2')(x)
x = Conv2D(512, (3, 3), activation='relu', padding='same', name='block5_conv3')(x)
x = MaxPooling2D((2, 2), strides=(2, 2), name='block5_pool')(x)
```

# VGG16 Implementation

```
# Block 4
x = Conv2D(512, (3, 3), activation='relu', padding='same', name='block4_conv1')(x)
x = Conv2D(512, (3, 3), activation='relu', padding='same', name='block4_conv2')(x)
x = Conv2D(512, (3, 3), activation='relu', padding='same', name='block4_conv3')(x)
x = MaxPooling2D((2, 2), strides=(2, 2), name='block4_pool')(x)

# Block 5
x = Conv2D(512, (3, 3), activation='relu', padding='same', name='block5_conv1')(x)
x = Conv2D(512, (3, 3), activation='relu', padding='same', name='block5_conv2')(x)
x = Conv2D(512, (3, 3), activation='relu', padding='same', name='block5_conv3')(x)
x = MaxPooling2D((2, 2), strides=(2, 2), name='block5_pool')(x)

# Flatten and dense layers
x = Flatten(name='flatten')(x)
x = Dense(4096, activation='relu', name='fc1')(x)
x = Dense(4096, activation='relu', name='fc2')(x)
output_tensor = Dense(num_classes, activation='softmax', name='predictions')(x)
```

# VGG16 Implementation

```
# Block 4
x = Conv2D(512, (3, 3), activation='relu', padding='same', name='block4_conv1')(x)
x = Conv2D(512, (3, 3), activation='relu', padding='same', name='block4_conv2')(x)
x = Conv2D(512, (3, 3), activation='relu', padding='same', name='block4_conv3')(x)
x = MaxPooling2D((2, 2), strides=(2, 2), name='block4_pool')(x)

# Block 5
x = Conv2D(512, (3, 3), activation='relu', padding='same', name='block5_conv1')(x)
x = Conv2D(512, (3, 3), activation='relu', padding='same', name='block5_conv2')(x)
x = Conv2D(512, (3, 3), activation='relu', padding='same', name='block5_conv3')(x)
x = MaxPooling2D((2, 2), strides=(2, 2), name='block5_pool')(x)

# Flatten and dense layers
x = Flatten(name='flatten')(x)
x = Dense(4096, activation='relu', name='fc1')(x)
x = Dense(4096, activation='relu', name='fc2')(x)
output_tensor = Dense(num_classes, activation='softmax', name='predictions')(x)

# Create model
model = Model(inputs=input_tensor, outputs=output_tensor, name='vgg16')

return model
```