

Deep Learning


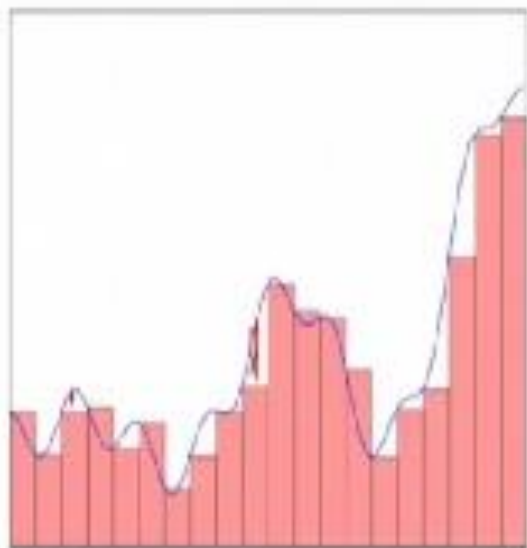
✧ Representative Power of Multilayer Networks

- A multilayer network of **perceptrons** with a single hidden layer can be used to approximate any Boolean function precisely
- A multilayer network of **sigmoid neurons** with a single hidden layer can be used to approximate any continuous function to any desired precision

↑ sigmoid

Multilayer Network

- For any function $f(x): \mathbb{R}^n \rightarrow \mathbb{R}^m$, we can find a network with enough neurons, whose output $g(x)$ satisfies $|g(x) - f(x)| < \epsilon$
- Such an arbitrary function can be represented by several tower functions




Multilayer Network

- All tower functions are similar and only differ in height and position on x-axis
- A black box takes some input and constructs a tower function
 - A network can add them up to approximate the function
- If we take the logistic function and set w to a very high value, we can recover step function
 - w controls the slope of the logistic function
- Can also adjust value of b to control position on x-axis at which function transitions from 0 to 1




Multilayer Network

Take two such sigmoid functions, with different b 's, and subtract them – will get a tower function




Multilayer Network

- More input parameters??
- Ex. 2 parameters



Single Hidden Layer Neural Network




Hidden Layer: States of nodes are unobserved


Inputs are densely connected to perceptrons, hence they are called **Dense** layers or **Fully Connected** layers

Feedforward Neural Network

- Input is an n -dimensional vector (0^{th} layer) $\in \mathbb{R}^n$
- Network has $L-1$ hidden layers
- 1 output layer containing k neurons (ex. for k classes)
- Each neuron – aggregation and activation




Feedforward Neural Network



Assuming n^i neurons in hidden layer h^i , $W^i \in \mathbb{R}^{n^{(i-1)} \times n^i}$ and $b^i \in \mathbb{R}^{n^i}$ between layers $i-1$ and i for $0 < i < L$


$W^L \in \mathbb{R}^{n^i \times k}$ and $b^L \in \mathbb{R}^k$ between last hidden layer and output layer



Aggregation at layer i : $\mathbf{z}^i = \mathbf{W}^i \mathbf{a}^{i-1} + \mathbf{b}^i$

For first hidden layer: $\mathbf{z}^1 = \mathbf{W}^1 \mathbf{a}^0 + \mathbf{b}^1$

$$\begin{pmatrix} z_1^1 \\ z_2^1 \\ z_3^1 \end{pmatrix} = \begin{pmatrix} W_{11} & W_{12} & W_{13} \\ W_{21} & W_{22} & W_{23} \\ W_{31} & W_{32} & W_{33} \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \\ x_3 \end{pmatrix} + \begin{pmatrix} b_1 \\ b_2 \\ b_3 \end{pmatrix} = \begin{pmatrix} \sum W_{1i} x_i + b_1 \\ \sum W_{2i} x_i + b_2 \\ \sum W_{3i} x_i + b_3 \end{pmatrix}$$



Activation at layer $i = g(\mathbf{z}^i) = g(\mathbf{b}^i + \mathbf{W}^i \mathbf{a}^{i-1})$

For first hidden layer: $g(\mathbf{z}^1) = g(\mathbf{b}^1 + \mathbf{W}^1 \mathbf{a}^0)$


$$\begin{bmatrix} a_1 \\ a_2 \\ a_3 \end{bmatrix} = \begin{bmatrix} g(z_1) \\ g(z_2) \\ g(z_3) \end{bmatrix}$$

Eg. $g(z_1) = \sigma(z_1) = 1 / (1 + e^{-z_1})$

g: activation function (logistic, tanh, linear etc.)

$$y = \tanh$$

$$y' = 1 - y^2$$



Aggregation at output layer $L = z^L = \mathbf{W}^L \mathbf{a}^{L-1} + \mathbf{b}^L$

$$z_1 = w_{11}x_1 + w_{12}x_2 + w_{13}x_3 + b$$

$$z_2 = w_{21}x_1 + w_{22}x_2 + w_{23}x_3 + b$$

Activation at output layer $L = \hat{\mathbf{y}} = g(z^L) = g(\mathbf{W}^L \mathbf{a}^{L-1} + \mathbf{b}^L)$

$$\begin{pmatrix} a_1 \\ a_2 \end{pmatrix} = \begin{pmatrix} g(z_1) \\ g(z_2) \end{pmatrix}$$

Learning parameters

In given example, dimensions of parameters:

- $\mathbf{W}^1: n^1 * n$ $\mathbf{b}^1: n^1$
- $\mathbf{W}^2: n^2 * n^1$ $\mathbf{b}^2: n^2$
- $\mathbf{W}^L: n^2 * k$ $\mathbf{b}^L: k$

- Assuming L layers and n^i neurons in hidden layer h^i and k neurons in output layer, no. of parameters to be learned:

- ✓ – Weights: $(L-1) * (n^{i-1} * n^i) + (n * k)$ for $0 < i < L$
- ✓ – Bias: $(L-1) * n^i + k$

Learning parameters

- **Data:** $\{x_i, y_i\}$ $i = 1..m$

- **Model:** $[w, b]$

$$\hat{y} = f(x) = g(W^3 g(W^2 g(W^1 x + b^1) + b^2) + b^3)$$

$$\hat{y} = [\hat{y}^1 \quad \hat{y}^2 \quad \dots \quad \hat{y}^k]$$

- **Algorithm:** Gradient Descent with back Propagation
- **Loss/Error function:** Sum of squared error loss

$$\min \frac{1}{N} \sum_{i=1}^m \sum_{j=1}^k (\hat{y}_j^i - y_j^i) \quad \text{for } i^{th} \text{ sample for all classes } j$$

]

Learning parameters

- Gradient Descent:

$t:=0;$

$max_iterations:=1000;$

Initialize $\theta_0 := [\mathbf{W}_0^1, \dots, \mathbf{W}_0^L, \mathbf{b}_0^1 \dots \mathbf{b}_0^L];$

while $t++ < max_iterations$ do

$$\theta_{t+1} := \theta_t - \eta \nabla \theta_t;$$

end

where, $\nabla \theta_t = \left[\frac{\partial L(\theta)}{\partial W_t}, \frac{\partial L(\theta)}{\partial b_t} \right]^T$

$\nabla \theta$ composed of:

– $\nabla W^1, \nabla W^2, \dots, \nabla W^{L-1} \in \mathbb{R}^{n(i-1) \times n_i}$, $\nabla W^L \in \mathbb{R}^{n \times k}$

– $\nabla b^1, \nabla b^2, \dots, \nabla b^{L-1} \in \mathbb{R}^{n_i}$, $\nabla b^L \in \mathbb{R}^k$

Loss function

- Loss function should capture how much \hat{y}_i deviates from y_i
- $y_i \in \mathbb{R}^n$ then squared error loss can be used:

$$L(\theta) = (1/m) * \sum (y_i - \hat{y}_i)^2$$

- Problems with squared error loss:

$$\frac{\partial L(w, b)}{\partial w} = (\hat{y} - y) * \hat{y} * (1 - \hat{y}) * x$$

– If $y_i = 1$ and $\hat{y}_i \sim 0$, $\frac{\partial L(w, b)}{\partial w} \sim 0$

Undesirable

– If $y_i = 0$ and $\hat{y}_i \sim 1$, $\frac{\partial L(w, b)}{\partial w} \sim 0$

Undesirable

– Weight updation becomes very slow

Loss function

- Information content (IC):
 - Events with high probability have low information content
 - “The sun will rise tomorrow”
 - Events with low probability have high information content
 - “There will be a cyclone tomorrow”
- $IC(A) = -\log_2(p(A))$
- Entropy: Expected information content = $\sum p_i * IC(i)$
 $= - \sum p_i \log_2(p_i)$

Loss function

Entropy: $y_i = [0 \quad 1 \quad 0 \quad 0]$

//Team B wins game

$\hat{y}_i = [0.2 \quad 0.1 \quad 0.4 \quad 0.3]$

//Our prediction

10K 5K 8K 1K

//Profit for each team win

Expected profit??

- Entropy: Expected information content = $\sum p_i IC(i)$

$$= - \sum p_i \log_2(p_i)$$

Actual predicted.

Loss function

- Cross-entropy: gives a measure on how close a predicted distribution is to a true distribution
 - True distribution p_i , Estimated distribution q_i
 - Estimated information content = $-\sum p_i \log_2(q_i)$
 - Capture difference between two probability distributions
 - If prediction is close to actual, cross entropy will be low

$$L(\theta) = -\sum y_c \log_2(\hat{y}_c)$$

$$y_c = 1$$

$$= 0$$

$$L(\theta) = -\log_2(\hat{y}_t)$$

for all k classes

if $c = t$ (true class)

otherwise

Loss function

- Objective function for classification:
 - Cross-entropy Loss

minimize: $L(\theta) = -\log_2(\hat{y}_t)$

$$\frac{\partial L}{\partial w} = -\frac{1}{\hat{y}_t} \frac{\partial \hat{y}_t}{\partial w}$$

\hat{y}_t : predicted probability of correct event


$\log_2(\hat{y}_t)$: probability that x belongs to t^{th} class, log-likelihood of data

Output Activation Function

- Output activation function:
 - Sum of outputs should be 1
 - \hat{y} should be a probability distribution
 - Sigmoid – probabilities will be $0 < p < 1$ but sum not equal to 1

$\sigma \rightarrow$
Softmax

$y_i = \{1 \quad 0 \quad 0 \quad 0\}$



$$\frac{e^{z_1}}{\sum_{i=1}^K e^{z_i}}$$

Classification problem

Output Activation Function

- Softmax function

$$z^L = b^L + W^L a^{L-1}$$

$$\hat{y} = g(z_j^L) = e^{z_j} / \sum e^{z_j}$$

for $j = 1..k$

z_j^L is j^{th} element of z^L



- Example: $z^L = [10 \quad 20 \quad -30]$

$$\hat{y} = [e^{10}/(e^{10} + e^{20} + e^{-30}) \quad e^{20}/(e^{10} + e^{20} + e^{-30}) \quad e^{-30}/(e^{10} + e^{20} + e^{-30})]$$

NOTE: Exponent converts -ve values to +ve values

Loss function

Chadwin




	Outputs	
	Real values	Probabilities
Output activation	Linear	Softmax
Loss function	Squared error	Cross-entropy

Backpropagation

How to compute $\nabla\theta$ composed of:

$$\nabla W^1, \nabla W^2, \dots, \nabla W^{L-1} \in \mathbb{R}^{n \times n}, \nabla W^L \in \mathbb{R}^{n \times k}$$

$$\nabla b^1, \nabla b^2, \dots, \nabla b^{L-1} \in \mathbb{R}^n, \nabla b^L \in \mathbb{R}^k$$




$$\frac{\partial L}{\partial w_{12}} = \frac{\partial L}{\partial \hat{y}} \times \frac{\partial \hat{y}}{\partial z^1} \times \frac{\partial z^1}{\partial a^1} \times \frac{\partial a^1}{\partial z^2} \times \frac{\partial z^2}{\partial a^2} \times \frac{\partial a^2}{\partial z^1} \times \frac{\partial z^1}{\partial w_{12}}$$

Backpropagation


Assuming classification problem, $L(\theta) = -\log_2(\hat{y}_t)$

- To learn weight w_{12}^1 use SGD and compute $\frac{\partial L(w,b)}{\partial W_{12}}$



Backpropagation

Assume a deep thin network, who is responsible for the loss??



Find derivative by chain rule:

$$\frac{\partial L(\theta)}{\partial W_{11}^1} = \underbrace{\frac{\partial L(\theta)}{\partial \hat{y}} * \frac{\partial \hat{y}}{\partial z_1^L}}_{\text{Output layer}} * \underbrace{\frac{\partial z_1^L}{\partial a_1^2} * \frac{\partial a_1^2}{\partial z_1^2}}_{\text{Previous hidden layer}} * \underbrace{\frac{\partial z_1^2}{\partial a_1^1} * \frac{\partial a_1^1}{\partial z_1^1}}_{\text{Previous hidden layer}} * \underbrace{\frac{\partial z_1^1}{\partial W_{11}^1}}_{\text{Weights}}$$


If we change W_{11} , how much does the loss change

Backpropagation

Assume a deep thin network

SSF

Cross Entropy = $-\sum y_i \log_2(\hat{y}_i)$



Find derivative by chain rule:

$$\frac{\partial L(\theta)}{\partial W_{11}^1} = \frac{\partial L(\theta)}{\partial \hat{y}} * \frac{\partial \hat{y}}{\partial z_1^L} * \frac{\partial z_1^L}{\partial a_1^2} * \frac{\partial a_1^2}{\partial z_1^2} * \frac{\partial z_1^2}{\partial a_1^1} * \frac{\partial a_1^1}{\partial z_1^1} * \frac{\partial z_1^1}{\partial W_{11}^1}$$

$$\frac{\partial L(\theta)}{\partial W_{11}^2} = \frac{\partial L(\theta)}{\partial \hat{y}} * \frac{\partial \hat{y}}{\partial z_1^L} * \frac{\partial z_1^L}{\partial a_1^2} * \frac{\partial a_1^2}{\partial z_1^2} * \frac{\partial z_1^2}{\partial W_{11}^2}$$

$$\frac{\partial L(\theta)}{\partial W_{11}^L} = \frac{\partial L(\theta)}{\partial \hat{y}} * \frac{\partial \hat{y}}{\partial z_1^L} * \frac{\partial z_1^L}{\partial W_{11}^L}$$


$$L(w, b) = -y \log \hat{y} - (1-y) \log (1-\hat{y})$$

$$\frac{\partial L}{\partial \hat{y}} = \left[\frac{-y}{\hat{y}} + \frac{(1-y)}{(1-\hat{y})} \right] \hat{y} (1-\hat{y}) = (\hat{y} - y)$$

$\frac{\partial L}{\partial \hat{y}} \times \frac{\partial \hat{y}}{\partial z}$

★ **BACKPROPAGATION WITH SIGMOID**
OUTPUT ACTIVATION & BINARY
CROSS-ENTROPY LOSS

Backpropagation



Assuming binary cross-entropy function

$$L = -y \log \hat{y} - (1 - y) \log (1 - \hat{y})$$


Loss \rightarrow Cross entropy

Activation \rightarrow sigmoid

$$\frac{\partial L(\theta)}{\partial \hat{y}} = \frac{-y}{\hat{y}} + \frac{1-y}{1-\hat{y}}$$

$$\frac{\partial L(\theta)}{\partial z} = \frac{\partial L(\theta)}{\partial \hat{y}} * \frac{\partial \hat{y}}{\partial z} = \left(\frac{-y}{\hat{y}} + \frac{1-y}{1-\hat{y}} \right) * \hat{y}(1-\hat{y}) = \hat{y} - y$$

$$\frac{\partial L(\theta)}{\partial w} = \frac{\partial L(\theta)}{\partial \hat{y}} * \frac{\partial \hat{y}}{\partial z} * \frac{\partial \hat{z}}{\partial w} = (\hat{y} - y) * x$$



$$\frac{\partial L(\theta)}{\partial \hat{y}} = \frac{-y}{\hat{y}} + \frac{1-y}{1-\hat{y}}$$

$$\frac{\partial L(\theta)}{\partial z^2} = \frac{\partial L(\theta)}{\partial \hat{y}} * \frac{\partial \hat{y}}{\partial z^2} = \left(\frac{-y}{\hat{y}} + \frac{1-y}{1-\hat{y}} \right) * \hat{y}(1-\hat{y}) = \hat{y} - y$$

$$\frac{\partial L(\theta)}{\partial w^2} = \frac{\partial L(\theta)}{\partial \hat{y}} * \frac{\partial \hat{y}}{\partial z^2} * \frac{\partial z^2}{\partial w^2} = (\hat{y} - y) * a^1$$

$$z^2 = w^2 a_1 + b$$

$$a^1 = \frac{1}{1 + e^{-z^1}}$$

$$:\delta^L = \frac{\partial L}{\partial \hat{y}} * \sigma'(z^L)$$

$$\hat{y} = \frac{1}{1 + e^{-z^2}}$$

$$z^1 = w_1 a_1 + b_1$$

$$\frac{\partial L(\theta)}{\partial a^1} = \frac{\partial L(\theta)}{\partial \hat{y}} * \frac{\partial \hat{y}}{\partial z^2} * \frac{\partial z^2}{\partial a^1} = (\hat{y} - y) * w^2$$

$$:\delta^L * w^2$$

$$\frac{\partial L(\theta)}{\partial z^1} = \frac{\partial L(\theta)}{\partial \hat{y}} * \frac{\partial \hat{y}}{\partial z^2} * \frac{\partial z^2}{\partial a^1} * \frac{\partial a^1}{\partial z^1} = (\hat{y} - y) * w^2 * a^1(1 - a^1)$$

$$:\delta^1 = \delta^L w^2 * \sigma'(z^1)$$

$$\frac{\partial L(\theta)}{\partial w^1} = \frac{\partial L(\theta)}{\partial \hat{y}} * \frac{\partial \hat{y}}{\partial z^2} * \frac{\partial z^2}{\partial a^1} * \frac{\partial a^1}{\partial z^1} * \frac{\partial z^1}{\partial w^1} = (\hat{y} - y) * w^2 * a^1(1 - a^1) * x$$


$$:\delta^1 * x$$

$$w^2$$

$$z^2 = w^2 a_1 + b^2$$

Backpropagation Equations

- $\delta^L = \frac{\partial L}{\partial \hat{y}} \odot \sigma'(z^L)$
- $\delta^l = ((w^{l+1})^T \delta^{l+1}) \odot \sigma'(z^l)$
- $\frac{\partial L}{\partial b_j^l} = \delta_j^l$
- $\frac{\partial L}{\partial w_{jk}^l} = a_k^{l-1} \delta_j^l$



Assuming:

Initial $w_{11}^2 = 12$, $w_{11}^1 = 0.1$

$\eta = 0.01$

During forward propagation:

$a_1^1 = 0.37$, $z_1^1 = 0.5$

$$\frac{\partial L}{\partial w_{11}^2} = \frac{\partial L}{\partial \hat{y}} * \frac{\partial \hat{y}}{\partial z_1^2} * \frac{\partial z_1^2}{\partial w_{11}^2}$$

$$\frac{\partial L}{\partial \hat{y}} * \frac{\partial \hat{y}}{\partial z_1^2} * \frac{\partial z_1^2}{\partial w_{11}^2}$$


$$\frac{\partial L(\theta)}{\partial \hat{y}} = \frac{-y}{\hat{y}} + \frac{1-y}{1-\hat{y}} = \frac{-1}{0.25} = -4$$

$$\frac{\partial L(\theta)}{\partial z_1^2} = \frac{\partial L(\theta)}{\partial \hat{y}} * \frac{\partial \hat{y}}{\partial z_1^2} = \hat{y} - y = 0.25 - 1 = -0.75$$

$$\left(\frac{-y}{\hat{y}} + \frac{1-y}{1-\hat{y}} \right) \hat{y}(1-\hat{y})$$

$$\frac{\partial L(\theta)}{\partial w_{11}^2} = \frac{\partial L(\theta)}{\partial \hat{y}} * \frac{\partial \hat{y}}{\partial z_1^2} * \frac{\partial z_1^2}{\partial w_{11}^2} = -0.75 * a_1^1 = -0.75 * 0.37 = -0.2775$$

$$w_{11}^{2*} = w_{11}^2 - \eta * \frac{\partial L}{\partial w_{11}^2} = 12 - 0.01 * (-0.2775) = 12.0028$$



$$\frac{\partial L}{\partial w_{11}^1} = \frac{\partial L}{\partial \hat{y}} * \frac{\partial \hat{y}}{\partial z_1^2} * \frac{\partial z_1^2}{\partial a_1^1} * \frac{\partial a_1^1}{\partial z_1^1} * \frac{\partial z_1^1}{\partial w_{11}^1}$$

$$= -0.75 * w_{11}^2 * a_1^1 (1 - a_1^1) * x_1$$

$$= -0.75 * 12 * 0.37 (1 - 0.37) * 60$$

$$= -125.874$$

$$w_{11}^{1*} = w_{11}^1 - \eta * \frac{\partial L}{\partial w_{11}^1} = 0.1 - 0.01 * (-125.874) = 1.35$$

Assignment: Compute $\frac{\partial L}{\partial w_{13}^1}$

BACKPROPAGATION WITH SOFTMAX OUTPUT ACTIVATION & CROSS-ENTROPY LOSS



Backpropagation

- Computing gradients
 - Gradients w.r.t. output units
 - Gradients w.r.t. hidden units
 - Gradients w.r.t. weights and biases

$$\frac{\partial L(\theta)}{\partial W_{11}^1} = \frac{\partial L(\theta)}{\partial \hat{y}} * \frac{\partial \hat{y}}{\partial z_1^L} * \frac{\partial z_1^L}{\partial a_1^2} * \frac{\partial a_1^2}{\partial z_1^2} * \frac{\partial z_1^2}{\partial a_1^1} * \frac{\partial a_1^1}{\partial z_1^1} * \frac{\partial z_1^1}{\partial W_{11}^1}$$

Gradients w.r.t. output units

- Assuming softmax activation and cross entropy loss at output layer for k classes:

$$L(\theta) = -\log_2(\hat{y}_t) \quad t: \text{true class label}$$

$$\frac{\partial L(\theta)}{\partial \hat{y}_i} = \frac{\partial(-\log \hat{y}_t)}{\partial \hat{y}_i} \quad \text{for } i = 1..k$$

$$= -\frac{1}{\hat{y}_t} \quad \text{if } i = t$$

$$= 0 \quad \text{otherwise}$$

Gradients w.r.t. output units

$$\begin{aligned}\frac{\partial L(\theta)}{\partial \hat{y}_i} &= -\frac{1}{\hat{y}_t} && \text{if } i = t \\ &= 0 && \text{otherwise}\end{aligned}$$


$$\begin{aligned}\frac{\partial L(\theta)}{\partial \hat{y}} &= \begin{pmatrix} \frac{\partial L(\theta)}{\partial \hat{y}_1} \\ \frac{\partial L(\theta)}{\partial \hat{y}_2} \\ \dots \\ \frac{\partial L(\theta)}{\partial \hat{y}_k} \end{pmatrix} = -\frac{1}{\hat{y}_t} * \begin{pmatrix} 1 \text{ if } t=1 \\ 1 \text{ if } t=2 \\ \dots \\ 1 \text{ if } t=k \end{pmatrix} \\ &= -\frac{1}{\hat{y}_t} * e(t)\end{aligned}$$

$e(t)$: One hot k -dimensional vector whose t^{th} entry is 1, others are 0

Gradients w.r.t. output units

$$\frac{\partial L(\theta)}{\partial z_i^L} = \frac{\partial(-\log \hat{y}_t)}{\partial z_i^L} = -\frac{\partial(-\log \hat{y}_t)}{\partial \hat{y}_t} * \frac{\partial \hat{y}_t}{\partial z_i^L} = -\frac{1}{\hat{y}_t} * \frac{\partial \hat{y}_t}{\partial z_i^L}$$

$$\hat{y}_t = \frac{\exp(z_t^L)}{\sum \exp(z_i^L)} \longrightarrow \hat{y}_t \text{ depends on } z_i^L$$



Gradients w.r.t. output units

$$\frac{\partial(-\log \hat{y}_t)}{\partial z_i^L} = -\frac{1}{\hat{y}_t} * \frac{\partial \hat{y}_t}{\partial z_i^L}$$

$$z^L = [z_1^L \quad z_2^L \quad \dots \quad z_k^L]$$

$$\hat{y} = \text{softmax}(z^L) = [\hat{y}_1^L \quad \hat{y}_2^L \quad \dots \quad \hat{y}_t^L \quad \dots \quad \hat{y}_k^L]$$

$$\hat{y}_t: t^{\text{th}} \text{ entry of } \hat{y} = \text{softmax}(z_t^L) = \frac{\exp(z_t^L)}{\sum \exp(z_i^L)}$$

$$\begin{aligned} \frac{\partial(-\log \hat{y}_t)}{\partial z_i^L} &= -\frac{1}{\hat{y}_t} * \frac{\partial \hat{y}_t}{\partial z_i^L} \\ &= -\frac{1}{\hat{y}_t} * \frac{\partial(\text{softmax}(z_t^L))}{\partial z_i^L} \end{aligned}$$

Gradients w.r.t. output units

$$\frac{\partial(-\log \hat{y}_t)}{\partial z_i^L} = -\frac{1}{\hat{y}_t} * \frac{\partial(\text{softmax}(z_t^L))}{\partial z_i^L} = -\frac{1}{\hat{y}_t} * \frac{\partial}{\partial z_i^L} \left(\frac{\exp(z_t^L)}{\sum \exp(z_i^L)} \right)$$

$$\begin{aligned} \frac{\partial(-\log \hat{y}_t)}{\partial z_i^L} &= -\frac{1}{\hat{y}_t} * \frac{\frac{\partial(\exp(z_t^L))}{\partial z_i^L}}{\sum \exp(z_i^L)} - \frac{\exp(z_t^L) * \frac{\partial(\sum \exp(z_i^L))}{\partial z_i^L}}{(\sum \exp(z_i^L))^2} \\ &= -\frac{1}{\hat{y}_t} * \left(\frac{1_{(t=i)} \exp(z_t^L)}{\sum \exp(z_i^L)} - \frac{\exp(z_t^L) * \exp(z_i^L)}{(\sum \exp(z_i^L))^2} \right) \\ &= -\frac{1}{\hat{y}_t} * (1_{(t=i)} \text{softmax}(z_t^L) - \text{softmax}(z_t^L) * \text{softmax}(z_i^L)) \\ &= -\frac{1}{\hat{y}_t} * (1_{(t=i)} \hat{y}_t - \hat{y}_t \hat{y}_i) = -(1_{(t=i)} - \hat{y}_i) \end{aligned}$$

$$\frac{\partial}{\partial x} \frac{f(x)}{g(x)} = \frac{g(x) * \frac{\partial(f(x))}{\partial x} - f(x) * \frac{\partial(g(x))}{\partial x}}{(g(x))^2} = \frac{\frac{\partial(f(x))}{\partial x}}{g(x)} - \frac{f(x) * \frac{\partial(g(x))}{\partial x}}{(g(x))^2}$$

Gradients w.r.t. output units

$$\frac{\partial L(\theta)}{\partial z_i^L} = - (1_{(t=i)} - \hat{y}_i)$$

Gradient w.r.t. vector \mathbf{z}^L :


$$\frac{\partial L(\theta)}{\partial \mathbf{z}^L} = \begin{pmatrix} \frac{\partial L(\theta)}{\partial z_1^L} \\ \frac{\partial L(\theta)}{\partial z_2^L} \\ \dots \\ \frac{\partial L(\theta)}{\partial z_k^L} \end{pmatrix} = \begin{pmatrix} - (1_{(t=1)} - \hat{y}_1) \\ - (1_{(t=2)} - \hat{y}_2) \\ \dots \\ - (1_{(t=k)} - \hat{y}_k) \end{pmatrix} = -(e(t) - \hat{\mathbf{y}})$$

$e(t)$: One hot k -dimensional vector whose t^{th} entry is 1, others are 0

Gradients w.r.t. hidden units

Example


$$\begin{aligned}\frac{\partial z}{\partial x} &= \frac{\partial z}{\partial y_1} * \frac{\partial y_1}{\partial x} + \frac{\partial z}{\partial y_2} * \frac{\partial y_2}{\partial x} \\ &= \sum \frac{\partial z}{\partial y_i} * \frac{\partial y_i}{\partial x} \quad \text{for } i = 1, 2\end{aligned}$$



Gradients w.r.t. hidden units


Example

$$\frac{\partial z}{\partial x} = \sum \frac{\partial z}{\partial y_i} * \frac{\partial y_i}{\partial x} \quad \text{for } i = 1, 2, 3$$



Chain rule across multiple paths


Gradients w.r.t. hidden units



$$z^{l+1} = W^{l+1} a_j^l + b^{l+1}$$

$$\frac{\partial L(\theta)}{\partial a_j^l} = \sum_{p=1}^k \frac{\partial L(\theta)}{\partial z_p^{l+1}} * \frac{\partial z_p^{l+1}}{\partial a_j^l}$$

k possible paths from output to hidden unit




$$z^{l+1} = W^{l+1} a_j^l + b^{l+1}$$

$$\begin{bmatrix} z_1^L \\ z_2^L \\ z_3^L \end{bmatrix} = \begin{bmatrix} w_{11}^L & w_{12}^L & w_{13}^L & w_{14}^L \\ w_{21}^L & w_{22}^L & w_{23}^L & w_{24}^L \\ w_{31}^L & w_{32}^L & w_{33}^L & w_{34}^L \end{bmatrix} \begin{bmatrix} a_1^2 \\ a_2^2 \\ a_3^2 \\ a_4^2 \end{bmatrix} + \begin{bmatrix} b_1^L \\ b_2^L \\ b_3^L \end{bmatrix}$$

$$z_1^L = W_{11}^L a_1^2 + W_{12}^L a_2^2 + W_{13}^L a_3^2 + W_{14}^L a_4^2 + b_1^L$$

$$\frac{\partial z_1^L}{\partial a_2^2} = w_{12}^L$$

Gradients w.r.t. hidden units



$$z^{l+1} = W^{l+1} a_j^l + b^{l+1}$$

$$\frac{\partial L(\theta)}{\partial a_j^l} = \sum_{p=1}^k \frac{\partial L(\theta)}{\partial z_p^{l+1}} * \frac{\partial z_p^{l+1}}{\partial a_j^l}$$

$$\frac{\partial L(\theta)}{\partial a_j^l} = \sum_{p=1}^k \frac{\partial L(\theta)}{\partial z_p^{l+1}} * W_{p,j}^{l+1}$$

rate

Gradients w.r.t. hidden units

$$\frac{\partial L(\theta)}{\partial a_j^l} = \sum_{p=1}^k \frac{\partial L(\theta)}{\partial z_p^{l+1}} * \frac{\partial z_p^{l+1}}{\partial a_j^l}$$

$$\frac{\partial L(\theta)}{\partial a_j^l} = \sum_{p=1}^k \frac{\partial L(\theta)}{\partial z_p^{l+1}} * W_{p,j}^{l+1}$$

$$\frac{\partial L(\theta)}{\partial z^{l+1}} = \begin{pmatrix} \frac{\partial L(\theta)}{\partial z_1^{l+1}} \\ \frac{\partial L(\theta)}{\partial z_2^{l+1}} \\ \dots \\ \frac{\partial L(\theta)}{\partial z_k^{l+1}} \end{pmatrix}$$

$$W_{\cdot,j}^{l+1} =$$

$$\begin{pmatrix} W_{1,j}^{l+1} \\ W_{2,j}^{l+1} \\ \dots \\ W_{k,j}^{l+1} \end{pmatrix}$$

① Vision

② Unsupervised

③ Adv. Analytic

$W_{\cdot,j}^{l+1}$ is the j^{th} column of W^{l+1}

$$(W_{\cdot,j}^{l+1})^T \frac{\partial L(\theta)}{\partial z^{l+1}} = \sum_{p=1}^k \frac{\partial L(\theta)}{\partial z_p^{l+1}} * W_{p,j}^{l+1}$$

Gradients w.r.t. hidden units

$$\frac{\partial L(\theta)}{\partial a_j^l} = (W_{\cdot,j}^{l+1})^\top \frac{\partial L(\theta)}{\partial z^{l+1}}$$

$$\frac{\partial L(\theta)}{\partial a_j^l} = \begin{pmatrix} \frac{\partial L(\theta)}{\partial a_1^l} \\ \frac{\partial L(\theta)}{\partial a_2^l} \\ \dots \\ \frac{\partial L(\theta)}{\partial a_n^l} \end{pmatrix} = \begin{pmatrix} (W_{\cdot,1}^{l+1})^\top \frac{\partial L(\theta)}{\partial z^{l+1}} \\ (W_{\cdot,2}^{l+1})^\top \frac{\partial L(\theta)}{\partial z^{l+1}} \\ \dots \\ (W_{\cdot,n}^{l+1})^\top \frac{\partial L(\theta)}{\partial z^{l+1}} \end{pmatrix}$$

$$\frac{\partial L(\theta)}{\partial a_j^l} = (W^{l+1})^\top \frac{\partial L(\theta)}{\partial z^{l+1}}$$


How to compute $\frac{\partial L(\theta)}{\partial z^{l+1}}$ for $l \leq L-1$

Gradients w.r.t. hidden units

$$\frac{\partial L(\theta)}{\partial z^l} = \begin{pmatrix} \frac{\partial L(\theta)}{\partial z_1^l} \\ \frac{\partial L(\theta)}{\partial z_2^l} \\ \dots \\ \frac{\partial L(\theta)}{\partial z_n^l} \end{pmatrix} \qquad \frac{\partial L(\theta)}{\partial z_j^l} = \frac{\partial L(\theta)}{\partial a_j^l} * \frac{\partial a_j^l}{\partial z_j^l} = \frac{\partial L(\theta)}{\partial a_j^l} * \frac{\partial g(z_j^l)}{\partial z_j^l}$$

$$\frac{\partial L(\theta)}{\partial z^l} = \begin{pmatrix} \frac{\partial L(\theta)}{\partial a_1^l} * \frac{\partial g(z_1^l)}{\partial z_1^l} \\ \frac{\partial L(\theta)}{\partial a_2^l} * \frac{\partial g(z_2^l)}{\partial z_2^l} \\ \dots \\ \frac{\partial L(\theta)}{\partial a_n^l} * \frac{\partial g(z_n^l)}{\partial z_n^l} \end{pmatrix} = \frac{\partial L(\theta)}{\partial a_j^l} \odot \left[\dots, \dots, \frac{\partial g(z_j^l)}{\partial z_j^l}, \dots \right]$$


Gradients w.r.t. Parameters



$$\frac{\partial L(\theta)}{\partial w^l} = \frac{\partial L(\theta)}{\partial z^l} * \frac{\partial z^l}{\partial w^l}$$

$$z^l = W^l a^{l-1} + b^l$$

Gradients w.r.t. Parameters



$$\frac{\partial L(\theta)}{\partial w^l} = \frac{\partial L(\theta)}{\partial z^l} * \frac{\partial z^l}{\partial w^l}$$

$$z^l = W^l a^{l-1} + b^l$$

$$\begin{pmatrix} z_1^l \\ z_2^l \\ \dots \\ z_n^l \end{pmatrix} = \begin{pmatrix} w_{11}^l & w_{12}^l & \dots & w_{1n}^l \\ w_{21}^l & w_{22}^l & \dots & w_{2n}^l \\ \dots & \dots & \dots & \dots \\ w_{n1}^l & w_{n2}^l & \dots & w_{nn}^l \end{pmatrix} \begin{pmatrix} a_1^{l-1} \\ a_2^{l-1} \\ \dots \\ a_n^{l-1} \end{pmatrix} + \begin{pmatrix} b_1^{l-1} \\ b_2^{l-1} \\ \dots \\ b_n^{l-1} \end{pmatrix}$$

Gradients w.r.t. Parameters

$$\frac{\partial L(\theta)}{\partial w^l} = \frac{\partial L(\theta)}{\partial z^l} * \frac{\partial z^l}{\partial w^l}$$

$$\begin{pmatrix} z_1^l \\ z_2^l \\ \dots \\ z_n^l \end{pmatrix} = \begin{pmatrix} w_{11}^l & w_{12}^l & \dots & w_{1n}^l \\ w_{21}^l & w_{22}^l & \dots & w_{2n}^l \\ \dots & \dots & \dots & \dots \\ w_{n1}^l & w_{n2}^l & \dots & w_{nn}^l \end{pmatrix} \begin{pmatrix} a_1^{l-1} \\ a_2^{l-1} \\ \dots \\ a_n^{l-1} \end{pmatrix} + \begin{pmatrix} b_1^{l-1} \\ b_2^{l-1} \\ \dots \\ b_n^{l-1} \end{pmatrix}$$

$$\frac{\partial z^l}{\partial w_{12}^l} = a_2^{l-1}$$

$$\frac{\partial z^l}{\partial w_{pq}^l} = a_q^{l-1}$$

$$\frac{\partial L(\theta)}{\partial w_{pq}^l} = \frac{\partial L(\theta)}{\partial z^l} * a_q^{l-1}$$

Gradients w.r.t. Parameters

$$\frac{\partial L(\theta)}{\partial w_{pq}^l} = \frac{\partial L(\theta)}{\partial z^l} * a_q^{l-1}$$

$$\frac{\partial L(\theta)}{\partial w^l} = \begin{pmatrix} \frac{\partial L(\theta)}{\partial w_{00}^l} & \frac{\partial L(\theta)}{\partial w_{01}^l} & \cdots & \frac{\partial L(\theta)}{\partial w_{0n-1}^l} \\ \cdots & \cdots & \cdots & \cdots \\ \cdots & \cdots & \cdots & \frac{\partial L(\theta)}{\partial w_{n-1\ n-1}^l} \end{pmatrix}$$

Backpropagation Equations

$$\frac{\partial L}{\partial z^l} = (w^{l+1})^T \cdot \frac{\partial L^l}{\partial z^{l+1}} * \sigma'(z^l)$$

$$\frac{\partial L}{\partial w^l} = \frac{\partial L^l}{\partial z^l} \cdot (a^{l-1})^T$$

$$\frac{\partial L}{\partial b^l} = \frac{\partial L^l}{\partial z^l}$$

* - element-wise multiplication

. - dot product

l – layer number

(assuming sigmoid activations)

Pseudo code: Gradient Descent

$t := 0;$

$max_iterations := 1000;$

Initialize $\theta_0 := [W_0^1, \dots, W_0^L, b_0^1 \dots b_0^L];$

while $t++ < max_iterations$ **do**

$a^1, a^2 \dots a^{L-1}, z^1, z^2 \dots z^L, \hat{y} = forward_propagation(\theta_t);$

$\nabla \theta_t = back_propagation(a^1, a^2 \dots a^{L-1}, z^1, z^2 \dots z^L, y, \hat{y});$

$\theta_{t+1} := \theta_t - \eta \nabla \theta_t;$

end

Pseudo code: Forward Propagation

for $v = 0$ to $L-1$ **do**

$$z^v = b^v + W^v a^{v-1};$$

$$a^v = g(z^v);$$

end;

$$z^L = b^L + W^L a^{L-1};$$

$$\hat{y} = O(z^L);$$

Do a forward propagation and compute all a^i 's, z^i 's, and \hat{y}

Pseudo code: Back Propagation

//Compute output gradient:

$$\nabla_z^L L(\theta) = - (e(y) - \hat{y});$$

for $v = L$ to 1 **do**

//Compute gradients w.r.t. parameters

$$\nabla_w^v L(\theta) = \nabla_z^v L(\theta) a^{v-1};$$

$$\nabla_b^v L(\theta) = \nabla_z^v L(\theta);$$

//Compute gradients w.r.t. layer below

$$\nabla_a^{v-1} L(\theta) = W^v \nabla_z^v L(\theta);$$

//Compute gradients w.r.t. layer below (pre-activation)

$$\nabla_z^{v-1} L(\theta) = \nabla_a^{v-1} L(\theta) \odot [\dots g'(z^{v-1,j}) \dots];$$

end