



Deep Learning

Representative Power of Multilayer Networks

- A multilayer network of **perceptrons** with a single hidden layer can be used to approximate any Boolean function precisely
- A multilayer network of **sigmoid neurons** with a single hidden layer can be used to approximate any continuous function to any desired precision

Multilayer Network

- For any function $f(x): \mathbb{R}^n \rightarrow \mathbb{R}^m$, we can find a network with enough neurons, whose output $g(x)$ satisfies $|g(x) - f(x)| < \epsilon$
- Such an arbitrary function can be represented by several tower functions




Multilayer Network

- All tower functions are similar and only differ in height and position on x-axis
- A black box takes some input and constructs a tower function
 - A network can add them up to approximate the function
- If we take the logistic function and set w to a very high value, we can recover step function
 - w controls the slope of the logistic function
- Can also adjust value of b to control position on x-axis at which function transitions from 0 to 1


Multilayer Network

Take two such sigmoid functions, with different b 's, and subtract them – will get a tower function




Multilayer Network

- More input parameters??
- Ex. 2 parameters



Single Hidden Layer Neural Network

2 Layer NN




Hidden Layer: States of nodes are unobserved


Inputs are densely connected to perceptrons, hence they are called **Dense** layers or **Fully Connected** layers

Feedforward Neural Network

- Input is an n -dimensional vector (0^{th} layer) $\in \mathbb{R}^n$
- Network has $L-1$ hidden layers
- 1 output layer containing k neurons (ex. for k classes)
- Each neuron – aggregation and activation




Feedforward Neural Network



Assuming n^i neurons in hidden layer h^i , $W^i \in \mathbb{R}^{n^{(i-1)} \times n^i}$ and $b^i \in \mathbb{R}^{n^i}$ between layers $i-1$ and i for $0 < i < L$


$W^L \in \mathbb{R}^{n^i \times k}$ and $b^L \in \mathbb{R}^k$ between last hidden layer and output layer



Aggregation at layer i : $\mathbf{z}^i = \mathbf{W}^i \mathbf{a}^{i-1} + \mathbf{b}^i$

For first hidden layer: $\mathbf{z}^1 = \mathbf{W}^1 \mathbf{a}^0 + \mathbf{b}^1$

$$\begin{pmatrix} z_1^1 \\ z_2^1 \\ z_3^1 \end{pmatrix} = \begin{pmatrix} W_{11} & W_{12} & W_{13} \\ W_{21} & W_{22} & W_{23} \\ W_{31} & W_{32} & W_{33} \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \\ x_3 \end{pmatrix} + \begin{pmatrix} b_1 \\ b_2 \\ b_3 \end{pmatrix} = \begin{pmatrix} \sum W_{1i} x_i + b_1 \\ \sum W_{2i} x_i + b_2 \\ \sum W_{3i} x_i + b_3 \end{pmatrix}$$




Activation at layer $i = g(\mathbf{z}^i) = g(\mathbf{b}^i + \mathbf{W}^i \mathbf{a}^{i-1})$
 For first hidden layer: $g(\mathbf{z}^1) = g(\mathbf{b}^1 + \mathbf{W}^1 \mathbf{a}^0)$

$$\begin{bmatrix} a_1 \\ a_2 \\ a_3 \end{bmatrix} = \begin{bmatrix} g(z_1) \\ g(z_2) \\ g(z_3) \end{bmatrix}$$

Eg. $g(z_1) = \sigma(z_1) = 1 / (1 + e^{-z_1})$

g: activation function (logistic, tanh, linear etc.)



Aggregation at output layer $L = z^L = \mathbf{W}^L \mathbf{a}^{L-1} + \mathbf{b}^L$

$$z_1 = w_{11}x_1 + w_{12}x_2 + w_{13}x_3 + b$$

$$z_2 = w_{21}x_1 + w_{22}x_2 + w_{23}x_3 + b$$

Activation at output layer $L = \hat{\mathbf{y}} = g(z^L) = g(\mathbf{W}^L \mathbf{a}^{L-1} + \mathbf{b}^L)$

$$\begin{bmatrix} a_1 \\ a_2 \end{bmatrix} = \begin{bmatrix} g(z_1) \\ g(z_2) \end{bmatrix}$$

Learning parameters

In given example, dimensions of parameters:

- $\mathbf{W}^1: n^1 * n$ $\mathbf{b}^1: n^1$
- $\mathbf{W}^2: n^2 * n^1$ $\mathbf{b}^2: n^2$
- $\mathbf{W}^L: n^2 * k$ $\mathbf{b}^L: k$
- Assuming L layers and n^i neurons in hidden layer h^i and k neurons in output layer, no. of parameters to be learned:
 - Weights: $(L-1) * (n^{i-1} * n^i) + (n * k)$ for $0 < i < L$
 - Bias: $(L-1) * n^i + k$

Learning parameters

- **Data:** $\{x_i, y_i\}$ $i = 1..m$

- **Model:**

$$\hat{y} = f(x) = g(W^3 g(W^2 g(W^1 x + b^1) + b^2) + b^3)$$

$$\hat{y} = [\hat{y}^1 \quad \hat{y}^2 \quad \dots \quad \hat{y}^k]$$

- **Algorithm:** Gradient Descent with back Propagation
- **Loss/Error function:** Sum of squared error loss

$$\min \frac{1}{N} \sum_{i=1}^m \sum_{j=1}^k (\hat{y}_j^i - y_j^i) \quad \text{for } i^{th} \text{ sample for all classes } j$$

Learning parameters

- Gradient Descent:

$t:=0;$

$max_iterations:=1000;$

Initialize $\theta_0 := [\mathbf{W}_0^1, \dots, \mathbf{W}_0^L, \mathbf{b}_0^1 \dots \mathbf{b}_0^L];$

while $t++ < max_iterations$ do

$\theta_{t+1} := \theta_t - \eta \nabla \theta_t;$

end

where, $\nabla \theta_t = \left[\frac{\partial L(\theta)}{\partial W_t}, \frac{\partial L(\theta)}{\partial b_t} \right]^T$

$\nabla \theta$ composed of:

– $\nabla W^1, \nabla W^2, \dots, \nabla W^{L-1} \in \mathbb{R}^{n(i-1) \times n_i}$, $\nabla W^L \in \mathbb{R}^{n \times k}$

– $\nabla b^1, \nabla b^2, \dots, \nabla b^{L-1} \in \mathbb{R}^{n_i}$, $\nabla b^L \in \mathbb{R}^k$

Loss function

- Loss function should capture how much \hat{y}_i deviates from y_i
- $y_i \in \mathbb{R}^n$ then squared error loss can be used:

$$L(\theta) = (1/N) * \sum (y_i - \hat{y}_i)^2$$

- Problems with squared error loss:

$$\frac{\partial L(w, b)}{\partial w} = (\hat{y} - y) * \hat{y} * (1 - \hat{y}) * x$$

- If $y_i = 1$ and $\hat{y}_i \sim 0$, $\frac{\partial L(w, b)}{\partial w} \sim 0$ Undesirable
- If $y_i = 0$ and $\hat{y}_i \sim 1$, $\frac{\partial L(w, b)}{\partial w} \sim 0$ Undesirable
- Weight updation becomes very slow

Loss function

- Information content (IC):
 - Events with high probability have low information content
 - “The sun will rise tomorrow”
 - Events with low probability have high information content
 - “There will be a cyclone tomorrow”
- $IC(A) = -\log_2(p(A))$
- Entropy: Expected information content = $\sum p_i * IC(i)$
 $= - \sum p_i \log_2(p_i)$

Loss function

Entropy: $y_i = [0$	1	0	0]	//Team B wins game
$\hat{y}_i = [0.2$	0.1	0.4	0.3]	//Our prediction
10K	5K	8K	1K	//Profit for each team win

Expected profit??

- Entropy: Expected information content = $\sum p_i IC(i)$
 $= - \sum p_i \log_2(p_i)$

Loss function

- Cross-entropy: gives a measure on how close a predicted distribution is to a true distribution
 - True distribution p_i , Estimated distribution q_i
 - Estimated information content = $-\sum p_i \log_2(q_i)$
 - Capture difference between two probability distributions
 - If prediction is close to actual, cross entropy will be low

$$L(\theta) = -\sum y_c \log_2(\hat{y}_c)$$

$$y_c = 1$$

$$= 0$$

for all k classes

if $c = t$ (true class)

otherwise

$$L(\theta) = -\log_2(\hat{y}_t)$$

Loss function

- Objective function for classification:
 - Cross-entropy Loss

minimize: $L(\theta) = -\log_2(\hat{y}_t)$


\hat{y}_t : predicted probability of correct event

$\log_2(\hat{y}_t)$: probability that x belongs to t^{th} class, log-likelihood of data

Output and Loss function

- Output activation function:
 - Sum of outputs should be 1
 - \hat{y} should be a probability distribution
 - Sigmoid – probabilities will be $0 < p < 1$ but sum not equal to 1

$$y_i = \{1 \quad 0 \quad 0 \quad 0\}$$



Classification problem

Output Activation Function

- Softmax function

$$z^L = b^L + W^L a^{L-1}$$

$$\hat{y} = g(z_j^L) = e^{z_j} / \sum e^{z_j} \quad \text{for } j = 1..k$$

z_j^L is j^{th} element of z^L

- Example: $z^L = [10 \quad 20 \quad -30]$

$$\hat{y} = [e^{10}/(e^{10} + e^{20} + e^{-30}) \quad e^{20}/(e^{10} + e^{20} + e^{-30}) \quad e^{-30}/(e^{10} + e^{20} + e^{-30})]$$

NOTE: Exponent converts –ve values to +ve values

Loss function


	Outputs	
	Real values	Probabilities
Output activation	Linear	Softmax
Loss function	Squared error	Cross-entropy

Backpropagation

How to compute $\nabla\theta$ composed of:

$$\nabla W^1, \nabla W^2, \dots, \nabla W^{L-1} \in \mathbb{R}^{n \times n}, \nabla W^L \in \mathbb{R}^{n \times k}$$


$$\nabla b^1, \nabla b^2, \dots, \nabla b^{L-1} \in \mathbb{R}^n, \nabla b^L \in \mathbb{R}^k$$



Backpropagation


Assuming classification problem, $L(\theta) = -\log_2(\hat{y}_t)$

- To learn weight w_{12}^1 use SGD and compute $\frac{\partial L(w,b)}{\partial w_{12}^1}$



Backpropagation

Assume a deep thin network, who is responsible for the loss??



Find derivative by chain rule:

$$\frac{\partial L(\theta)}{\partial W_{11}^1} = \underbrace{\frac{\partial L(\theta)}{\partial \hat{y}} * \frac{\partial \hat{y}}{\partial z_1^L}}_{\text{Output layer}} * \underbrace{\frac{\partial z_1^L}{\partial a_1^2} * \frac{\partial a_1^2}{\partial z_1^2}}_{\text{Previous hidden layer}} * \underbrace{\frac{\partial z_1^2}{\partial a_1^1} * \frac{\partial a_1^1}{\partial z_1^1}}_{\text{Previous hidden layer}} * \underbrace{\frac{\partial z_1^1}{\partial W_{11}^1}}_{\text{Weights}}$$

If we change W_{11} , how much does the loss change

$$L = -y \log \hat{y} - (1-y) \log (1-\hat{y})$$

$$\frac{\partial L}{\partial \hat{y}} = \frac{-y}{\hat{y}} + \frac{1-y}{1-\hat{y}}$$

$$\begin{aligned} \frac{\partial \hat{y}}{\partial z} &= \sigma(z)(1-\sigma(z)) \\ &= \hat{y}(1-\hat{y}) \end{aligned}$$

$$\begin{aligned} \frac{\partial L}{\partial z} &= \frac{\partial L}{\partial \hat{y}} \times \frac{\partial \hat{y}}{\partial z} \\ &= \hat{y}(1-\hat{y}) \left(\frac{-y}{\hat{y}} + \frac{1-y}{1-\hat{y}} \right) \end{aligned}$$

$$= -y(1-\hat{y}) + \hat{y}(1-y)$$

$$= -y$$

$$= \hat{y} - y$$

$$\frac{\partial L}{\partial w} = \frac{\partial L}{\partial y} \times \frac{\partial y}{\partial z} \times \frac{\partial z}{\partial w}$$

$$= (\hat{y} - y) \alpha.$$

$$\frac{\partial L}{\partial a'} = \frac{\partial L}{\partial \hat{y}} \times \frac{\partial \hat{y}}{\partial z^2} \times \frac{\partial z^2}{\partial a'}$$

$$\Rightarrow z^2 = \omega_2 a_1 + b$$

$$= (\hat{y} - y) \times \omega^2$$

$$\frac{\partial L}{\partial z'} = \frac{\partial L}{\partial \hat{y}} \times \frac{\partial \hat{y}}{\partial z^2} \times \frac{\partial z^2}{\partial a_1} \times \frac{\partial a_1}{\partial z'}$$

$$\frac{\partial L}{\partial w_{11}^2} = \underbrace{\frac{\partial L}{\partial \hat{y}} \times \frac{\partial \hat{y}}{\partial z}} \times \frac{\partial z}{\partial w}$$

$$\frac{\partial L}{\partial \hat{y}} = \frac{-y}{\hat{y}} + \frac{1-y}{1-\hat{y}} = -4$$

$$= -4$$

$$\frac{\partial L}{\partial z} = \frac{\partial L}{\partial \hat{y}} \times \frac{\partial \hat{y}}{\partial z} = \hat{y} - y = -0.75$$

$$\frac{\partial L}{\partial w} = (-0.75)(0.37) = -0.2775$$

$$w_{11}^2 = w_{11}^2 - \eta \times \frac{\partial L}{\partial w_{11}^2} = 12 - 0.00(-0.2775)$$

$$= 12.0028$$

$$\frac{\partial L}{\partial w_{11}'} = \frac{\partial L}{\partial \hat{y}} \times \frac{\partial \hat{y}}{\partial z_1^2} \times \frac{\partial z_1^2}{\partial a_1'} \times \frac{\partial a_1'}{\partial w_{11}'} \quad ||$$

$$\frac{\partial a_1'}{\partial z_1'} \times \frac{\partial z_1'}{\partial w_{11}'}$$

$$= -0.75 \times 12 \times 0.5 (1-0.5)^{60}$$

$$= -13.5$$

Homework

$$\frac{\partial L}{\partial w_{13}'} =$$